

# z/OS 2.4 IBM Education Assistant (IEA)

Solution (Epic) Name: XCF Transport Class Simplification

Element(s)/Component(s): z/OS XCF



# Agenda

- Trademarks
- Session Objectives
- Overview
- Usage & Invocation
- Interactions & Dependencies
- Migration & Coexistence Considerations
- Installation
- Session Summary
- Appendix

# Trademarks

- See url <http://www.ibm.com/legal/copytrade.shtml> for a list of trademarks.
- Additional Trademarks:
  - None.

# Session Objectives

- A little background on XCF Transport Classes to provide some context
- Where are we headed?
- What do you get in z/OS V2R4?
  - How enabled?
  - What changes?
  - Caveats and concerns?
  - Validation?

We're not there yet.

But progress is being made.

# Background

Brief review of Transport Class basics

# Background: Transport Classes

- Transport classes segregate message traffic by size, group, or both
- Message traffic is segregated by size for the purpose of:
  - Ensuring timely transfer for small messages
  - Efficient utilization of buffer space
- Message traffic is segregated by group for the purpose of:
  - Isolating an ill-behaved group to protect others from harm
  - Dedicating XCF signal resources to a favored group so that it need not compete with others. Ostensibly, the group will:
    - Never be denied service
    - Never experience transfer delays

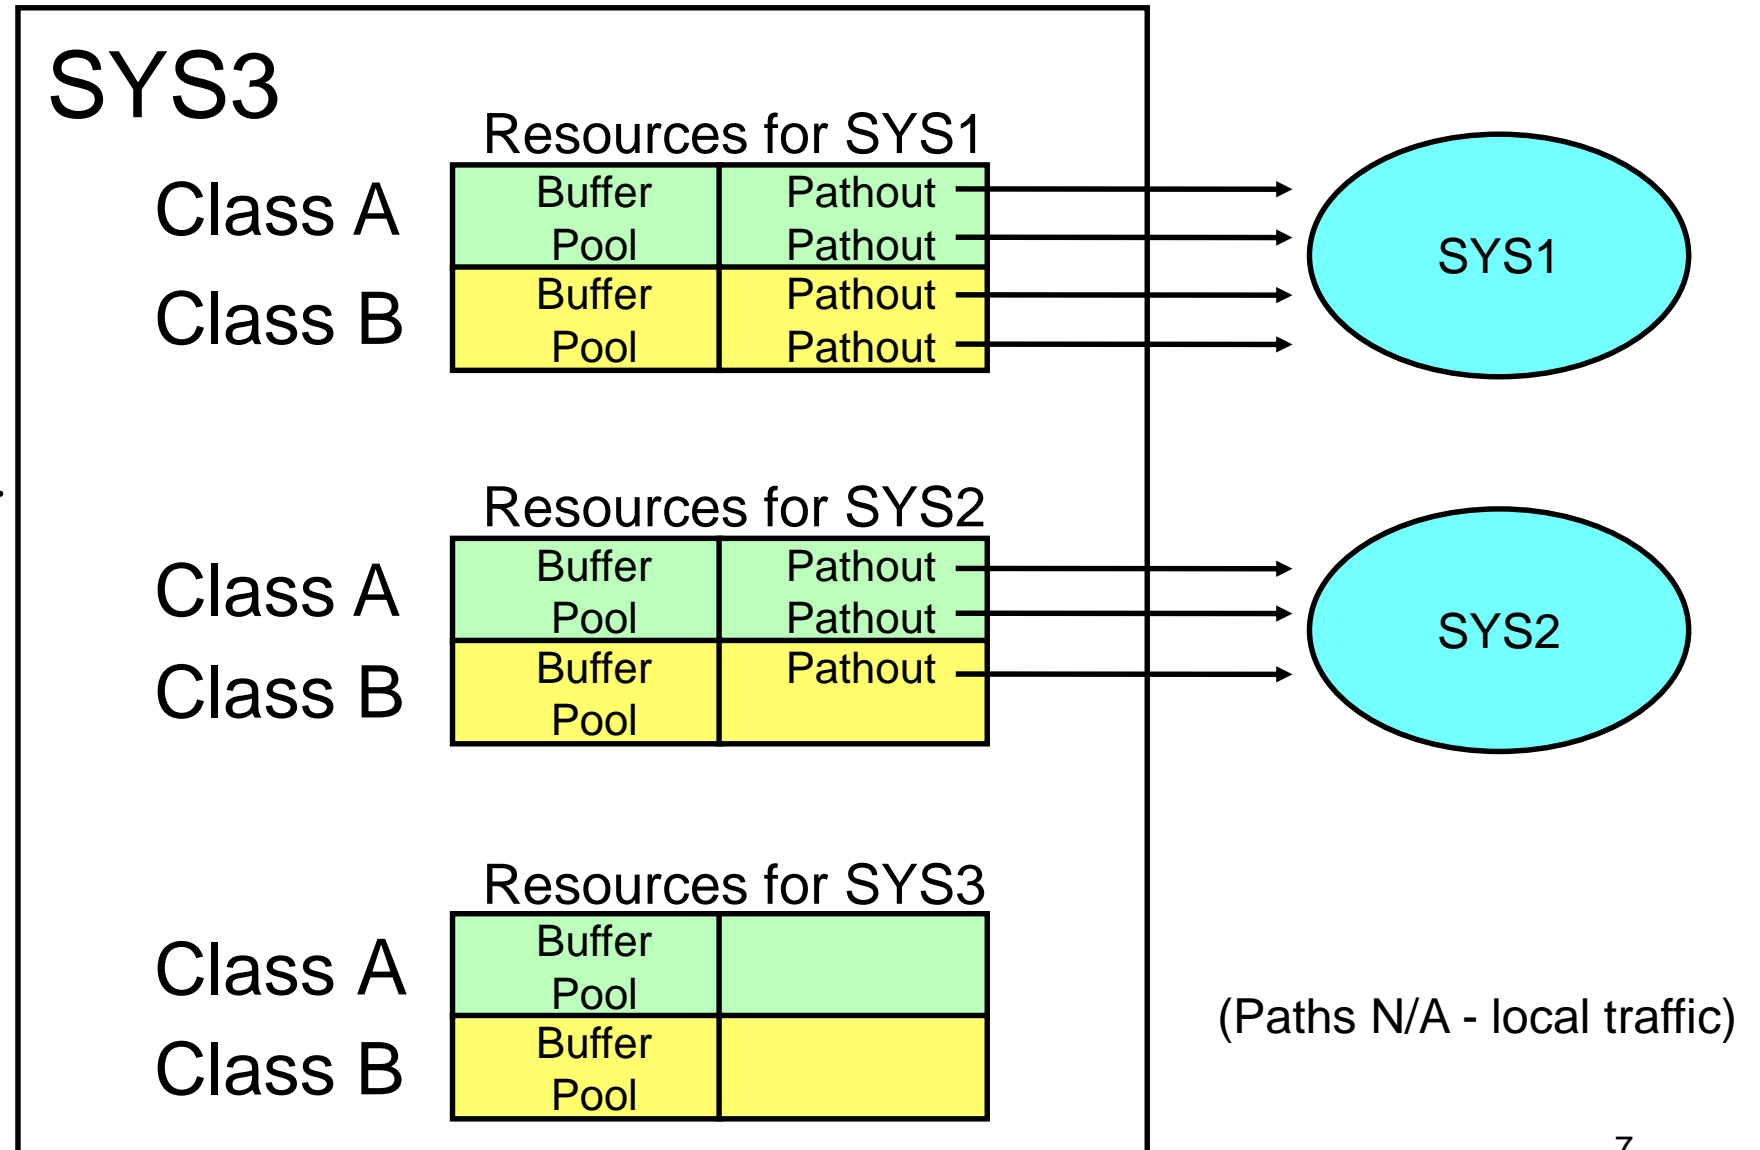
# Background: Transport Classes

Transport class has buffers and signal paths (unless local).

Each target system has it's own dedicated signal resources. Never shared between systems.

The same class definition applies to all target systems.

In practice, the signal traffic patterns are often vary dramatically on a target system basis.



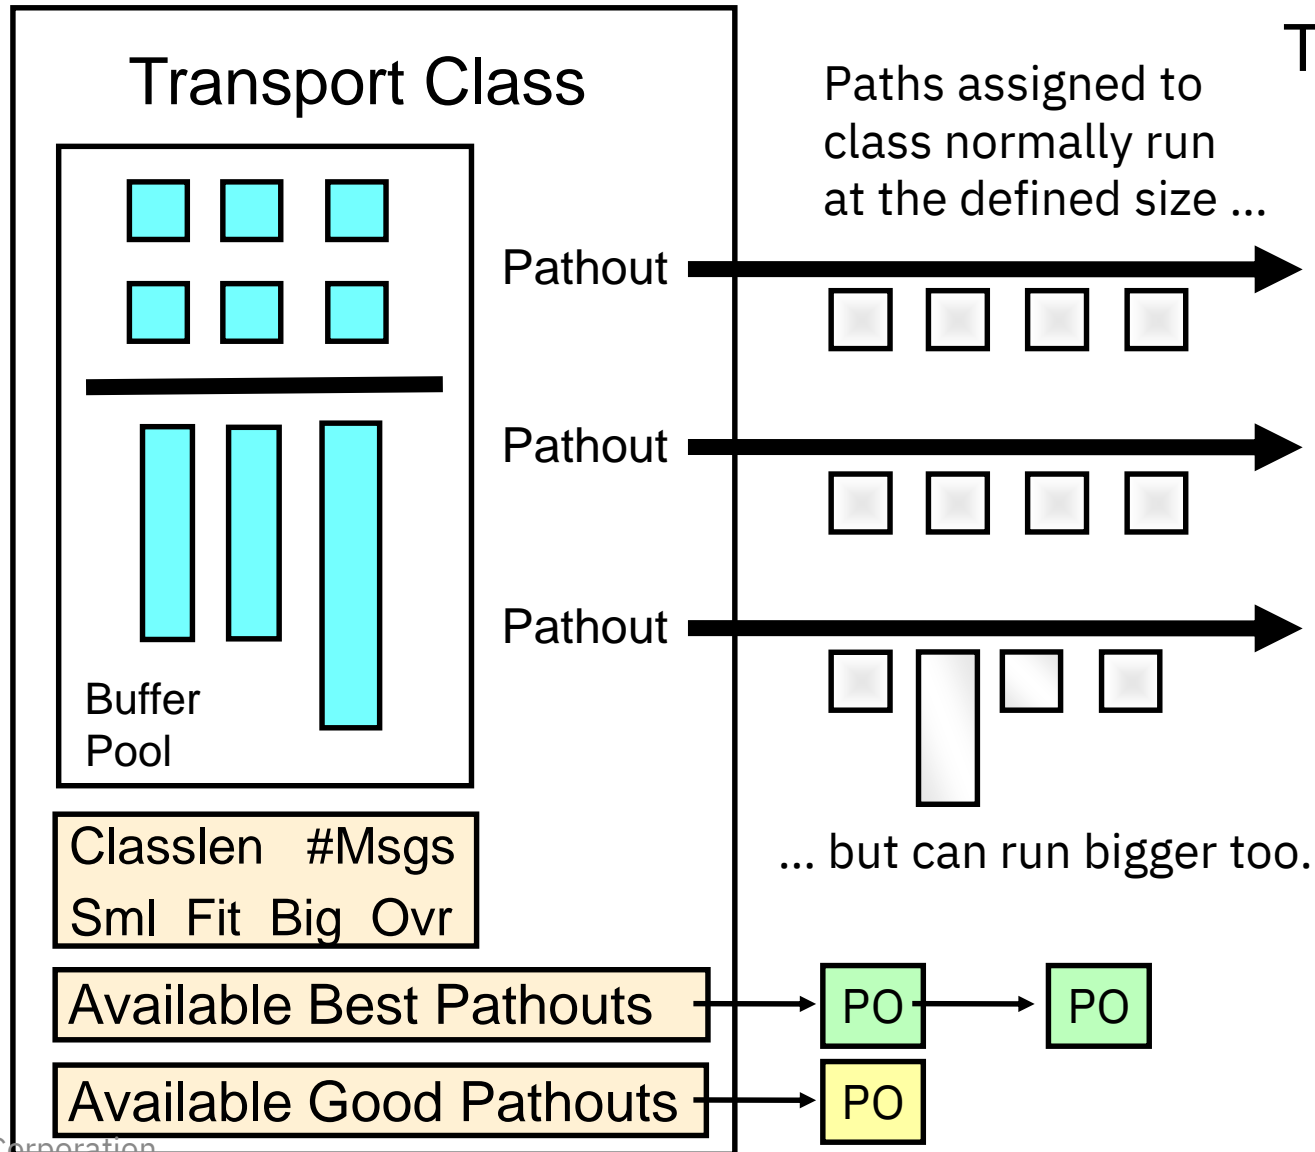
# Background: Transport Class

Buffers for signals no bigger than the “defined” size.


Buffers for signals longer than defined size.

Buffer Accounting

Path Accounting  
(Avail vs Busy)



## Target System



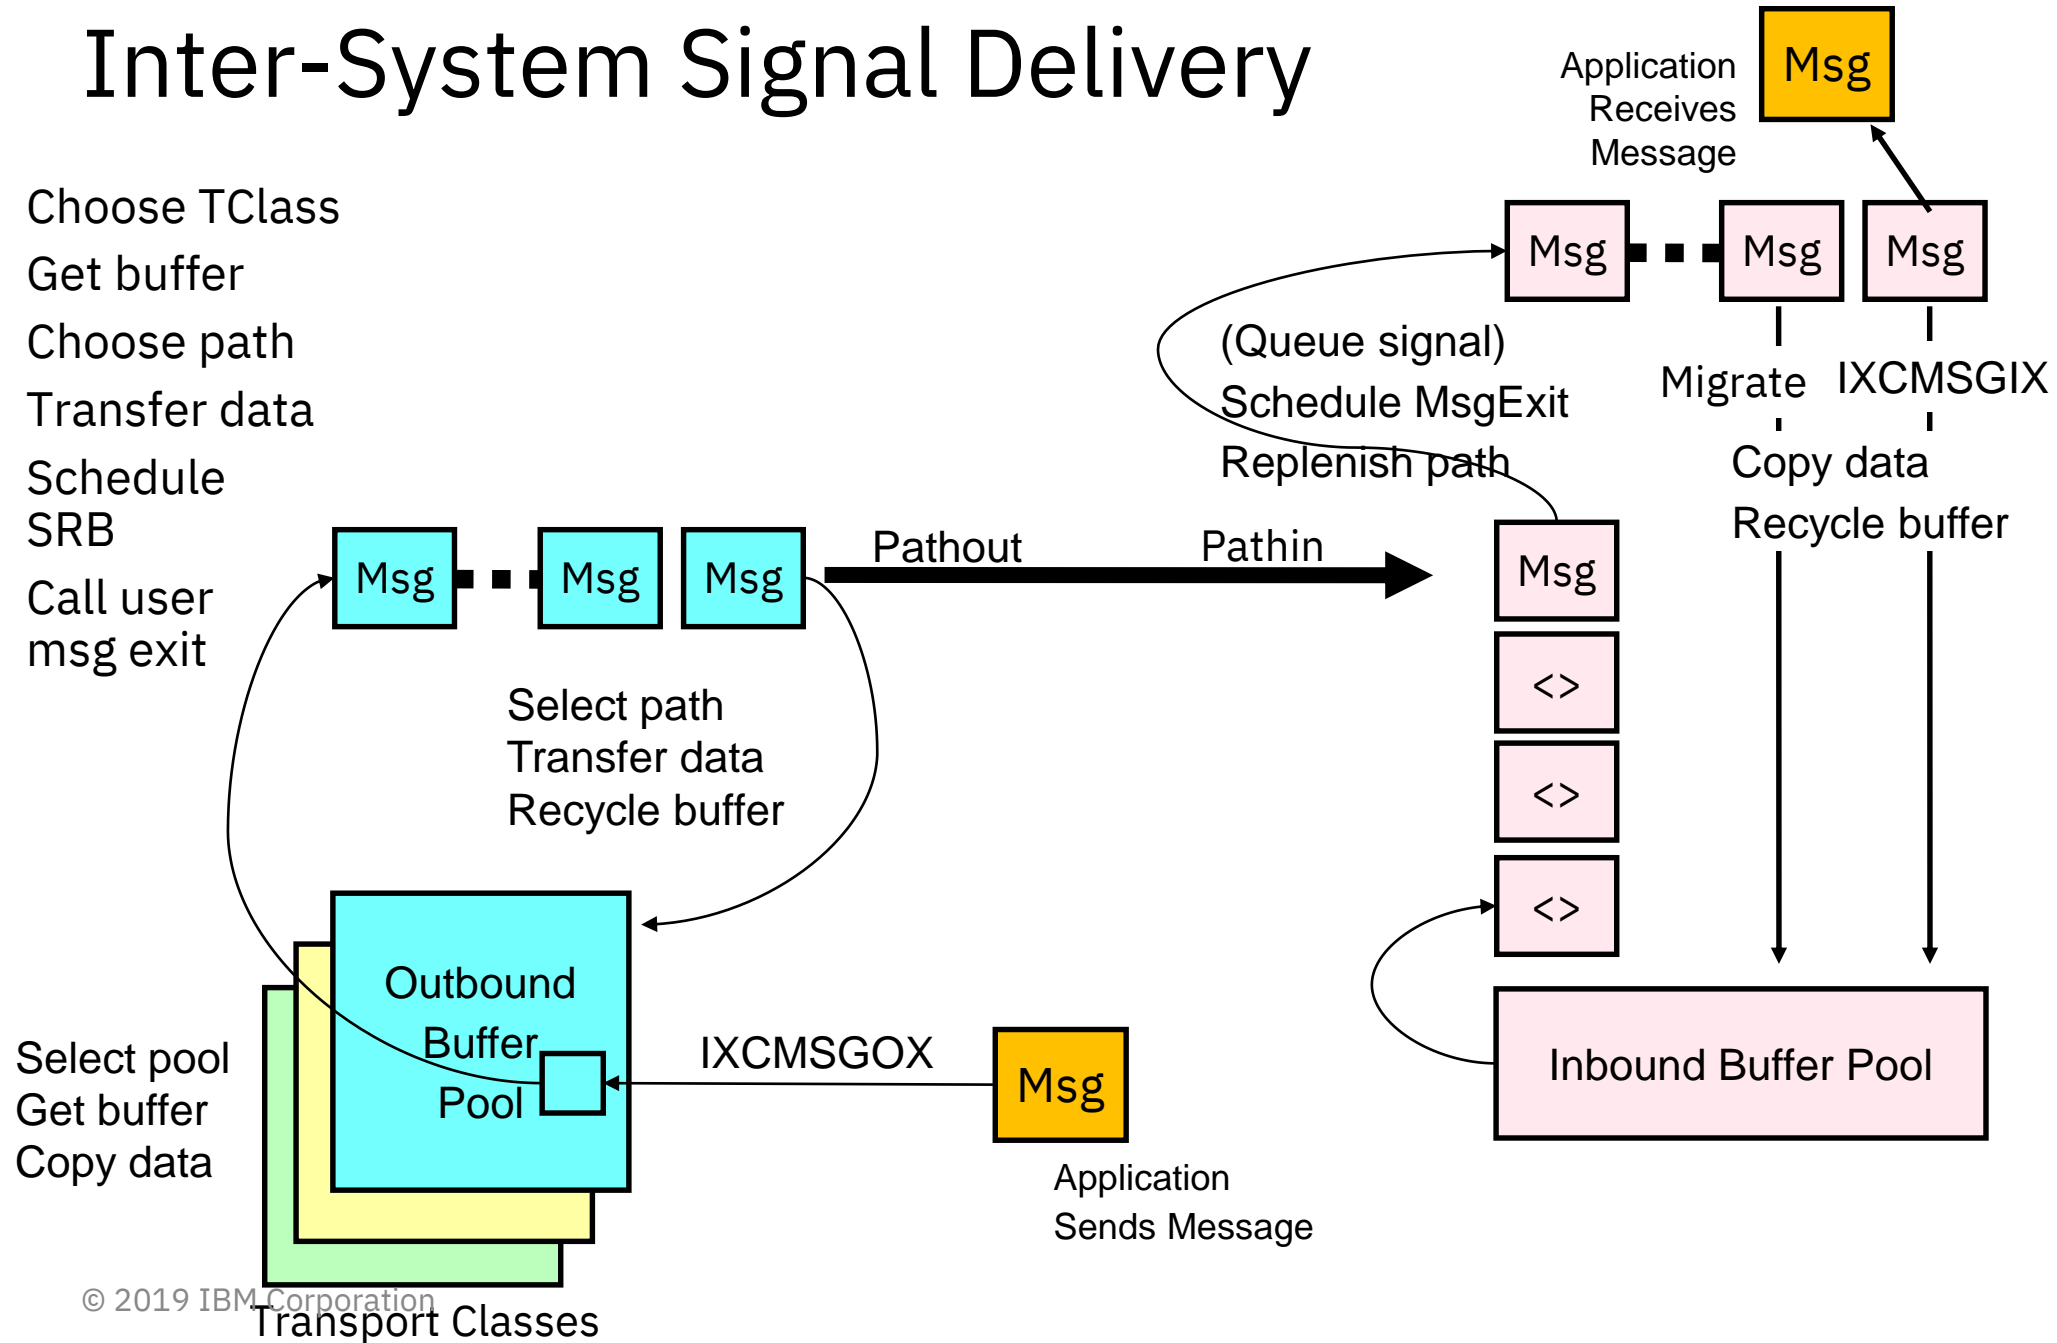
Inbound paths must be ready to receive whatever size signal is demanded by the outbound side.

Changes must be renegotiated through a process that sends additional control signals and temporarily quiesces signal transfer.



# Inter-System Signal Delivery

- Choose TClass
- Get buffer
- Choose path
- Transfer data
- Schedule SRB
- Call user msg exit



# Where are we headed?

The installation can achieve the high level of resiliency, availability, and performance expected of sysplex without having to plan, define, monitor, or tune z/OS XCF transport classes.

# Why eliminate the need to define Transport Classes?

- Simplification
  - Remove this burden from customers and all the IBM personnel that have to analyze, explain, make recommendations, etc.
  - RMF reports do not always provide clear guidance as to how signal resources should be configured; more art than science
- Avoid Outages
  - Transport Classes often not well understood, mistakes are made that lessen the resiliency of the sysplex and so permit avoidable outages
  - A Transport Class can help isolate an ill behaved member, but not really practical since you have to know in advance who is going to cause trouble
    - Such as delays and overwhelming bursts
- Self-optimization
  - Static definitions are not well suited to the dynamics of the sysplex
    - Can lead to inefficient use of resources (idle paths, excess storage)
  - System should automatically apply resources where needed most
    - For example, to handle a burst of messages or a change in signal patterns

# To eliminate the need to define Transport Classes

## **XCF must automatically handle the problems that transport classes were intended to address:**

- Timely message transfer
    - Maintain signal throughput and minimize signal delivery times, especially for the small messages that are typically most predominant in the sysplex
    - Provides better resiliency (minimize delays and queueing)
  - Efficient utilization of signal resources
    - Helps minimize cost
    - Provides better resiliency (more capacity)
- 
- Isolation of ill behaved members
    - Avoid sympathy sickness so that problems with signal delivery for one member don't negatively impact signal delivery for other members
  - Fair access to signal resources
    - Don't allow one member to monopolize the signal resources to the detriment of signal delivery for other members

Size segregation

Group segregation

# To eliminate the need to define Transport Classes

**XCF must automatically handle the problems that transport classes were intended to address:**

- Timely message transfer
  - Maintain signal throughput and minimize signal delivery times, especially for the small messages that are typically most predominant in the sysplex
  - Provides better resiliency

In the early days of sysplex, there was a significant difference in transfer time for small vs large signals. Today, this difference is much, much smaller and signal throughput is much greater. The relative delay arising from a small signal following a large one vs that of following a small signal is negligible. Intermixing signal sizes on a signal path will not have any discernable impact on the performance of the workload.

signal delivery for other members

For z/OS V2R4, timely message transfer has been a primary focus. Our rationale:

- There will likely be some consolidation of signal paths as transport classes are eliminated, we want to make sure the survivors can sustain the signal load.
- What if there is a workload for which those marginal differences in transfer times for small vs large signals does matter? The ability to sustain higher signal rates will likely mask any impacts that might arise.
- And ....

# To eliminate the need to define Transport Classes

## XCF must automatically handle the problems that transport classes were intended to address:

- Timely message transfer
  - Maintain signal throughput and minimize signal delay
  - Small messages that are typically most prevalent
  - Provides better resiliency
- Efficient utilization of signal resources
  - Helps minimize cost
  - Provides better resiliency
- Isolation of ill behaved members
  - Avoid sympathy sickness so that problems don't negatively impact signal delivery for other members
- Fair access to signal resources
  - Don't allow one member to monopolize the signal delivery for other members

These days, the more compelling reason for segregating signals by size has been to ensure that we can **maximize the number of signal buffers** available to the inbound path within the constraint of its MAXMSG limit. In general, more buffers tends to improve resiliency and helps maintain throughput.

Since our solution for eliminating the need to segregate signals by size uses the maximum size inbound buffer, we effectively **minimize the number of signal buffers** available to the inbound path. Even more impetus to focus on timely transfer and signal throughput!

# To eliminate the need to define Transport Classes

## **XCF must automatically handle the problems that transport classes were intended to address:**

- Timely message transfer
  - Maintain signal throughput and minimize signal delivery times, especially for the small messages that are typically most prevalent
  - Provides better resiliency
- Efficient utilization of signal resources
  - Helps minimize cost
  - Provides better resiliency
- Isolation of ill behaved members
  - Avoid sympathy sickness so that problem members don't negatively impact signal delivery for other members
- Fair access to signal resources
  - Don't allow one member to monopolize the signal delivery for other members

“Message Isolation” was delivered in z/OS V2R2. It is a foundational, incremental step along the journey towards the goal of eliminating the need to define transport classes for the purpose of segregating signals by group. Might provide “good enough” isolation, but it is not sufficient to make the claim that there is no need to define transport classes to segregate signals by group.

# To eliminate the need to define Transport Classes

## **XCF must automatically handle the problems that transport classes were intended to address:**

- Timely message transfer
  - Maintain signal throughput and minimize signal delivery times, especially for the small messages that are typically most problematic
  - Provides better resiliency
- Efficient utilization of signal resources
  - Helps minimize cost
  - Provides better resiliency
- Isolation of ill behaved members
  - Avoid sympathy sickness so that problem members don't negatively impact signal delivery for other members
- Fair access to signal resources
  - Don't allow one member to monopolize the signal delivery for other members

In general, we don't see field problems that would have been avoided if only we had had fair access to the XCF signal service. What I think of as "isolation" issues are more prevalent and problematic. But until we can guarantee fairness, we can't claim there is no need to define transport classes to segregate signals by group.



# So what do you get?

Elimination of the need for transport classes to segregate signals by size.

*(Elimination of need for group segregation is deferred).*

# Overview

- Who (Audience)
    - System programmers, sysplex architects, performance analysts, capacity analysts, diagnosticians
  - What (Solution)
    - XCF will internally manage the signal resources to provide timely delivery of signals in a sysplex independent of any Transport Class definitions **created purely for size segregation**.
  - Wow (Benefit / Value, Need Addressed)
    - Simplification: You no longer need to define, monitor, tune, or manage XCF Transport Class definitions to segregate signals **purely by size**.
    - Simplification: You need only configure an appropriate number of signal paths.
    - Resilience: Less potential for non-optimal transport class definitions to negatively impact signal delivery. Probably more signal capacity.<sup>1</sup>
- <sup>1</sup> Given the same configuration.

# The solution

- XCF will intermix signal sizes on “XCF Managed” signal paths as it sees fit while maintaining signal throughput and timely signal transfer, especially for small signals
  - Thus eliminating the need for transport class based signal size segregation
  - Any available path can be used for any size signal
    - Could do this today if we like. The challenge is to achieve this without incurring any discernable impact to signal throughput, system overhead, resiliency, ...

***We assume adequate signal capacity and system performance.***

***If inadequate capacity, or systems not performing well:***

- Transport classes are irrelevant; they cannot address such issues
- Segregating signal traffic on the sending system does not resolve issues related to the target system being unresponsive

# Feedback during development

- On the use of existing transport class definitions:
  - Either honor them as given, or ignore them outright
  - Want to control whether system uses new or old behavior
- Should be able to observe the new behavior
  - Want to be able to see which behavior is in play, new or old
- Focus first on eliminating the need for “size segregation”
  - Applicable to most if not all installations
  - For those that don’t do group segregation, the “problem” is solved
- Don’t change meaning of existing measurements (SMF data)
  - Differences between releases makes it hard to interpret the data

← So there is a new switch, XTCSIZE

← So there is a new XCF defined pseudo transport class, \_XCFMGD

# Usage & Invocation

- New XCF FUNCTIONS switch XTCSIZE determines what transport class segregation rules are available to XCF
- When XTCSIZE is DISABLED:
  - XCF signal resources will be managed per traditional transport class segregation rules
- When XTCSIZE is ENABLED:
  - XCF has the option to manage signal resources for **selected** transport classes per new “XCF Managed” segregation rules
  - Transport Classes subject to being XCF Managed per XTCSIZE are those that are defined purely for size segregation

# Interactions & Dependencies

- Exploitation requires:
  - A z/OS V2R4 sending system with XTCSIZE switch ENABLED
  - A target system running z/OS V2R4
    - Includes self, but the more interesting case is when target is some other system
- Software Dependencies
  - None.
- Hardware Dependencies
  - None.
- Exploiters
  - Implicit (potentially any user of XCF Signal Service)

# Migration & Coexistence Considerations

- You will always see the new pseudo `_XCFMGD` “transport class”
  - But does not impact system behavior unless `XTCSIZE` is `ENABLED`.
- No toleration/coexistence APARs/PTFs are needed.
  - z/OS V2R4 manages signal resources per the traditional transport class segregation rules when communicating with a system running an older z/OS release (the `XTCSIZE` switch setting is irrelevant).
  - So until all systems in the sysplex are running z/OS V2R4, you will still need your traditional transport class definitions.
- Migration actions:
  - For now, `XTCSIZE` is `ENABLED` by default. So ...
  - If you don't want the new behavior, `XTCSIZE` must be `DISABLED`:
    - `COUPLExx` parmlib member: `FUNCTIONS DISABLE(XTCSIZE)`
    - Operator command: `SETXCF FUNCTIONS,DISABLE=XTCSIZE`

# Installation

- IPL z/OS V2R4
  - With the code you are getting now, XTCSIZE is ENABLED by default.
    - Might be DISABLED for code delivered at general availability.
  - So for now, you will get new behavior upon IPL.
- Note that XTCSIZE is a “local” switch.
  - A system obeys its own local switch setting
  - If ENABLED, there is the potential for new behavior when sending signals to any system in the sysplex that is also running z/OS V2R4
  - The XTCSIZE switch setting on some other system in the sysplex does not influence local system behavior at all.
  - Purely a “send side” setting.



# Which classes are “XCF Managed” per XTCSIZE?

```
CLASSDEF CLASS(SML) CLASSLEN(956) GROUP(UNDESIG)
CLASSDEF CLASS(MED) CLASSLEN(12000)
CLASSDEF CLASS(BIG) CLASSLEN(32000)

CLASSDEF CLASS(BADGUY0) CLASSLEN(956) GROUP(MYAPP1)
CLASSDEF CLASS(BADGUY1) CLASSLEN(4028) GROUP(MYAPP1,UNDESIG)
CLASSDEF CLASS(BADGUY2) CLASSLEN(8124) GROUP(UNDESIG)
```

**XTCSIZE selects those transport class definitions that do NOT have a group explicitly assigned to the class.**

Which classes have explicitly assigned groups? These will NOT be XCF Managed. All others will be.

Classes BADGUY0 and BADGUY1 have an explicit group assignment, namely group MYAPP1.

GROUP(UNDESIG) is NOT an explicit group assignment. UNDESIG refers to all groups NOT explicitly assigned.

**Answer: XTCSIZE would manage/select these classes:**

DEFAULT – this class is implicitly defined with GROUP(UNDESIG) if you don’t otherwise change it.

SML - UNDESIG is not an explicit group assignment.

MED - no explicit groups assigned, so pure size.

BIG - no explicit groups assigned, so pure size.

BADGUY2 – UNDESIG is not an explicit group assignment.

BADGUY0 and BADGUY1 do suggest a desire for size segregation, but it’s not “pure size”. They are for group segregation.

# A subtle point on class selection

```
CLASSDEF CLASS(SML) CLASSLEN(956) GROUP(UNDESIG)
CLASSDEF CLASS(MED) CLASSLEN(12000)
CLASSDEF CLASS(BIG) CLASSLEN(32000)

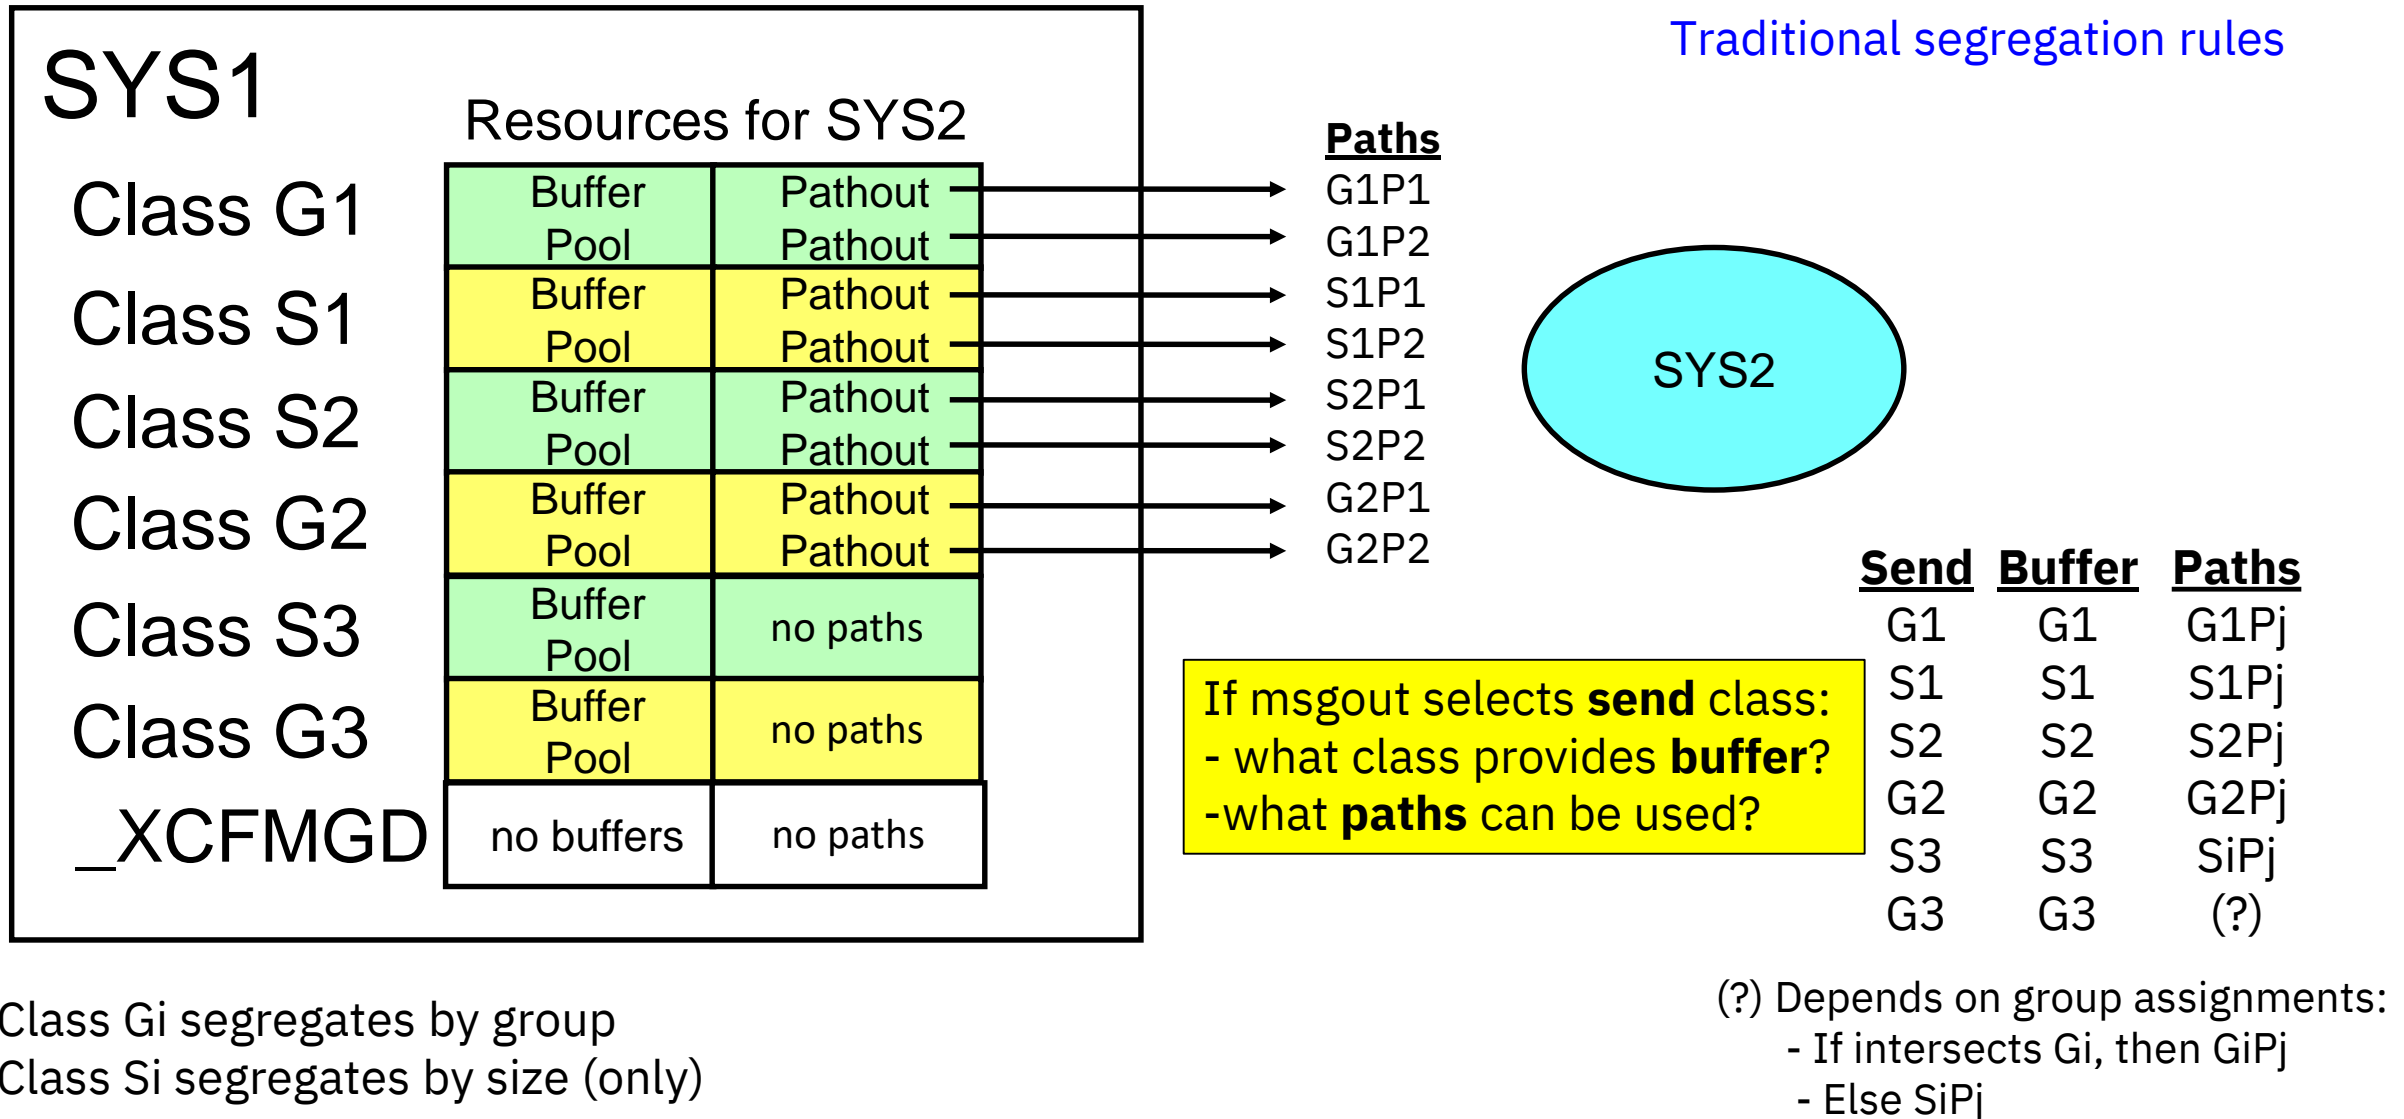
CLASSDEF CLASS(BADGUY0) CLASSLEN(956) GROUP(MYAPP1)
CLASSDEF CLASS(BADGUY1) CLASSLEN(4028) GROUP(MYAPP1,UNDESIG)
CLASSDEF CLASS(BADGUY2) CLASSLEN(8124) GROUP(UNDESIG)
```

Class BADGUY1 is actually available for use by any group, since MYAPP1+UNDESIG = everyone. Assume XTCSIZE is ENABLED.

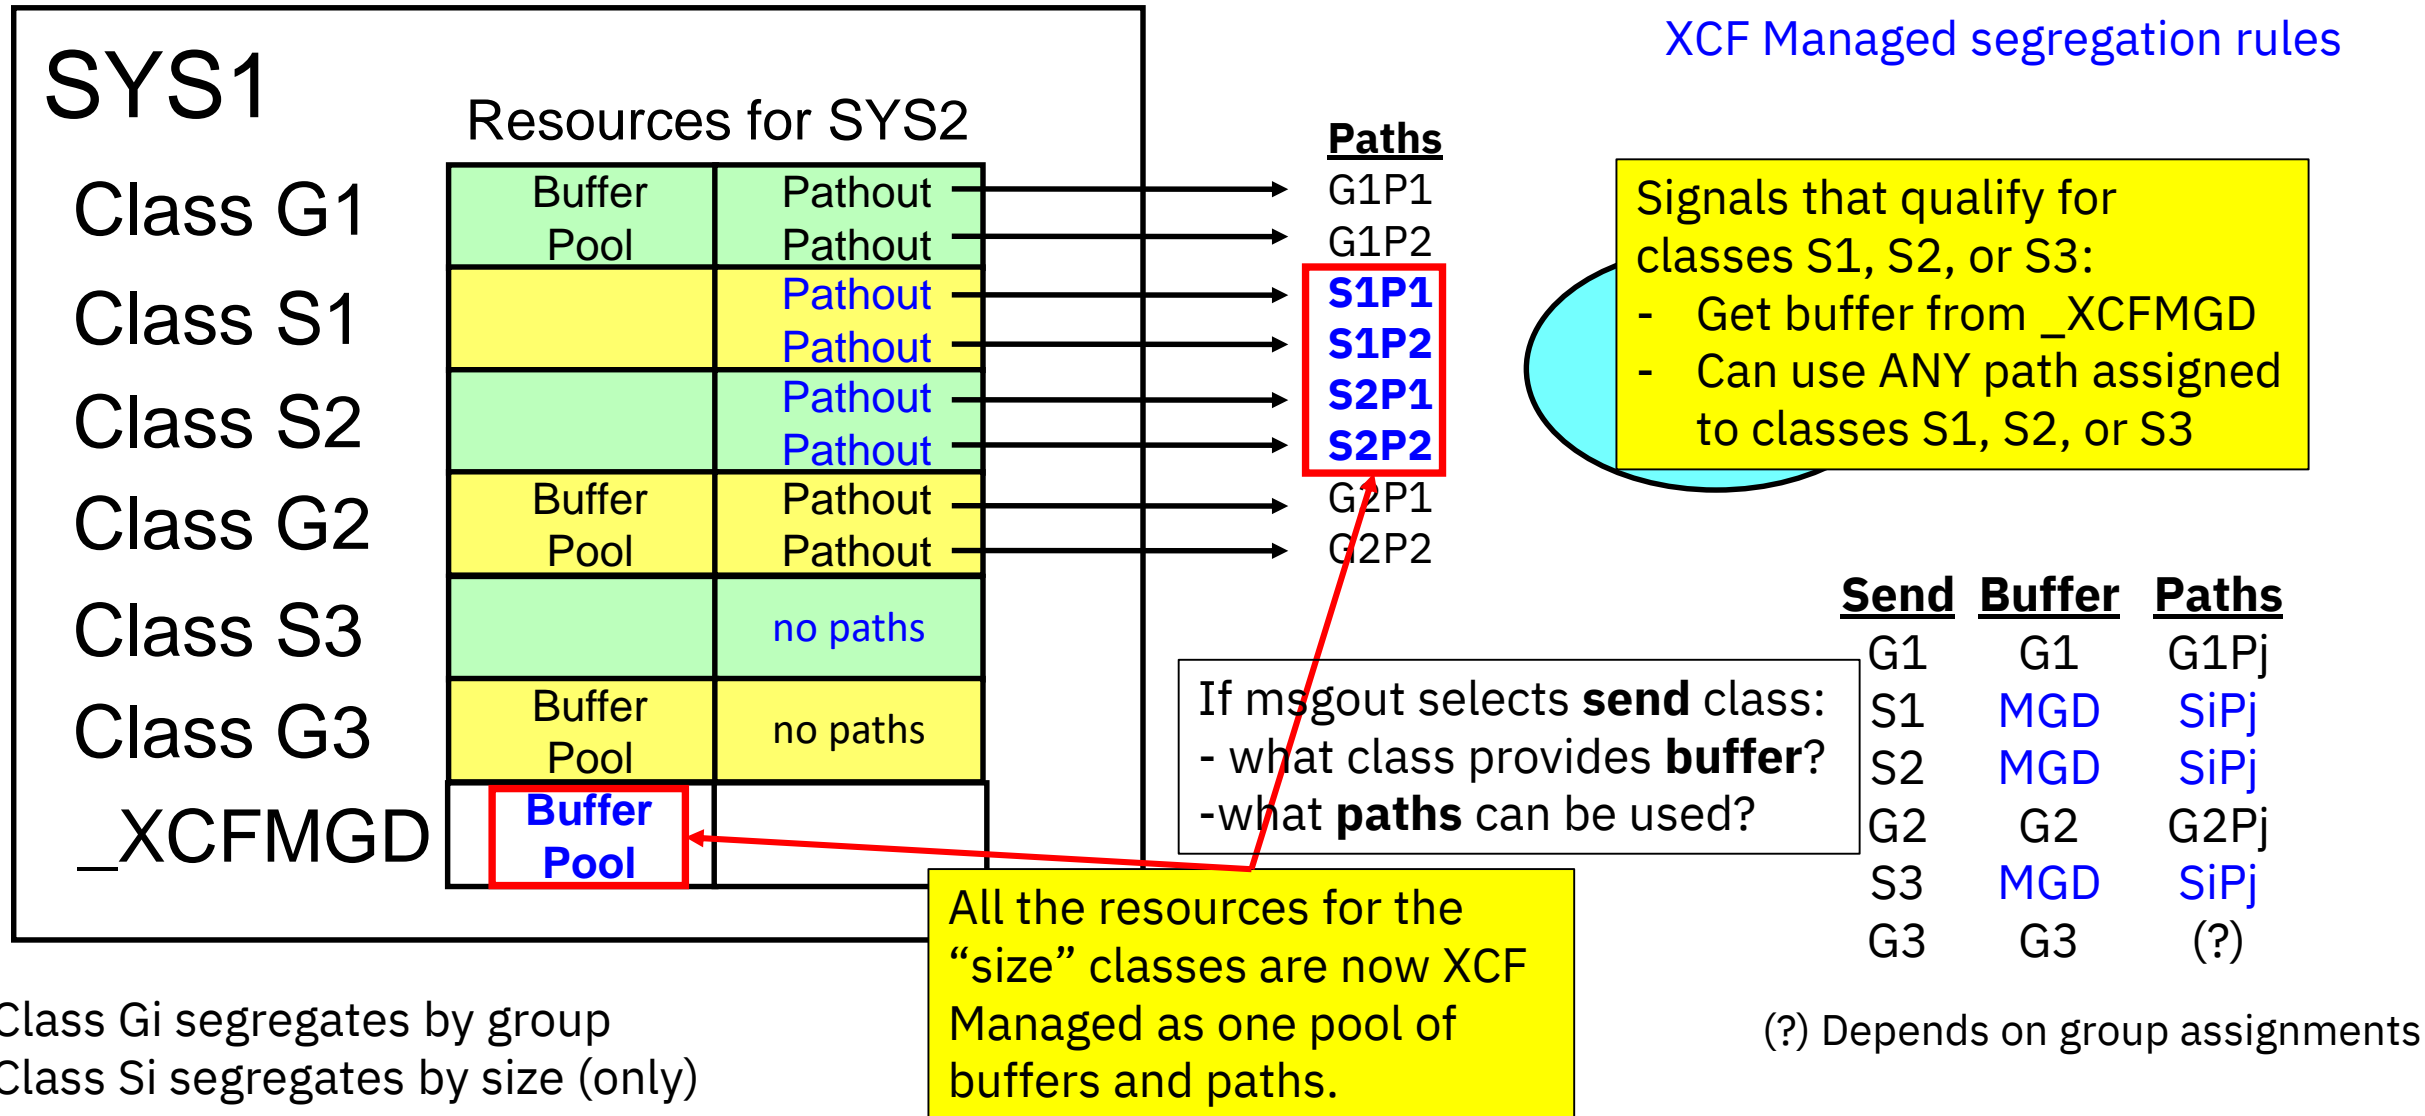
When an UNDESIG group sends a signal, both \_XCFMGD and BADGUY1 are candidate classes. BADGUY1 could be selected, and likely will be if BADGUY1 has paths and CLASSLEN(4028) is a best fit.

**So don't make the mistake of thinking that all UNDESIG signals must flow through \_XCFMGD.**

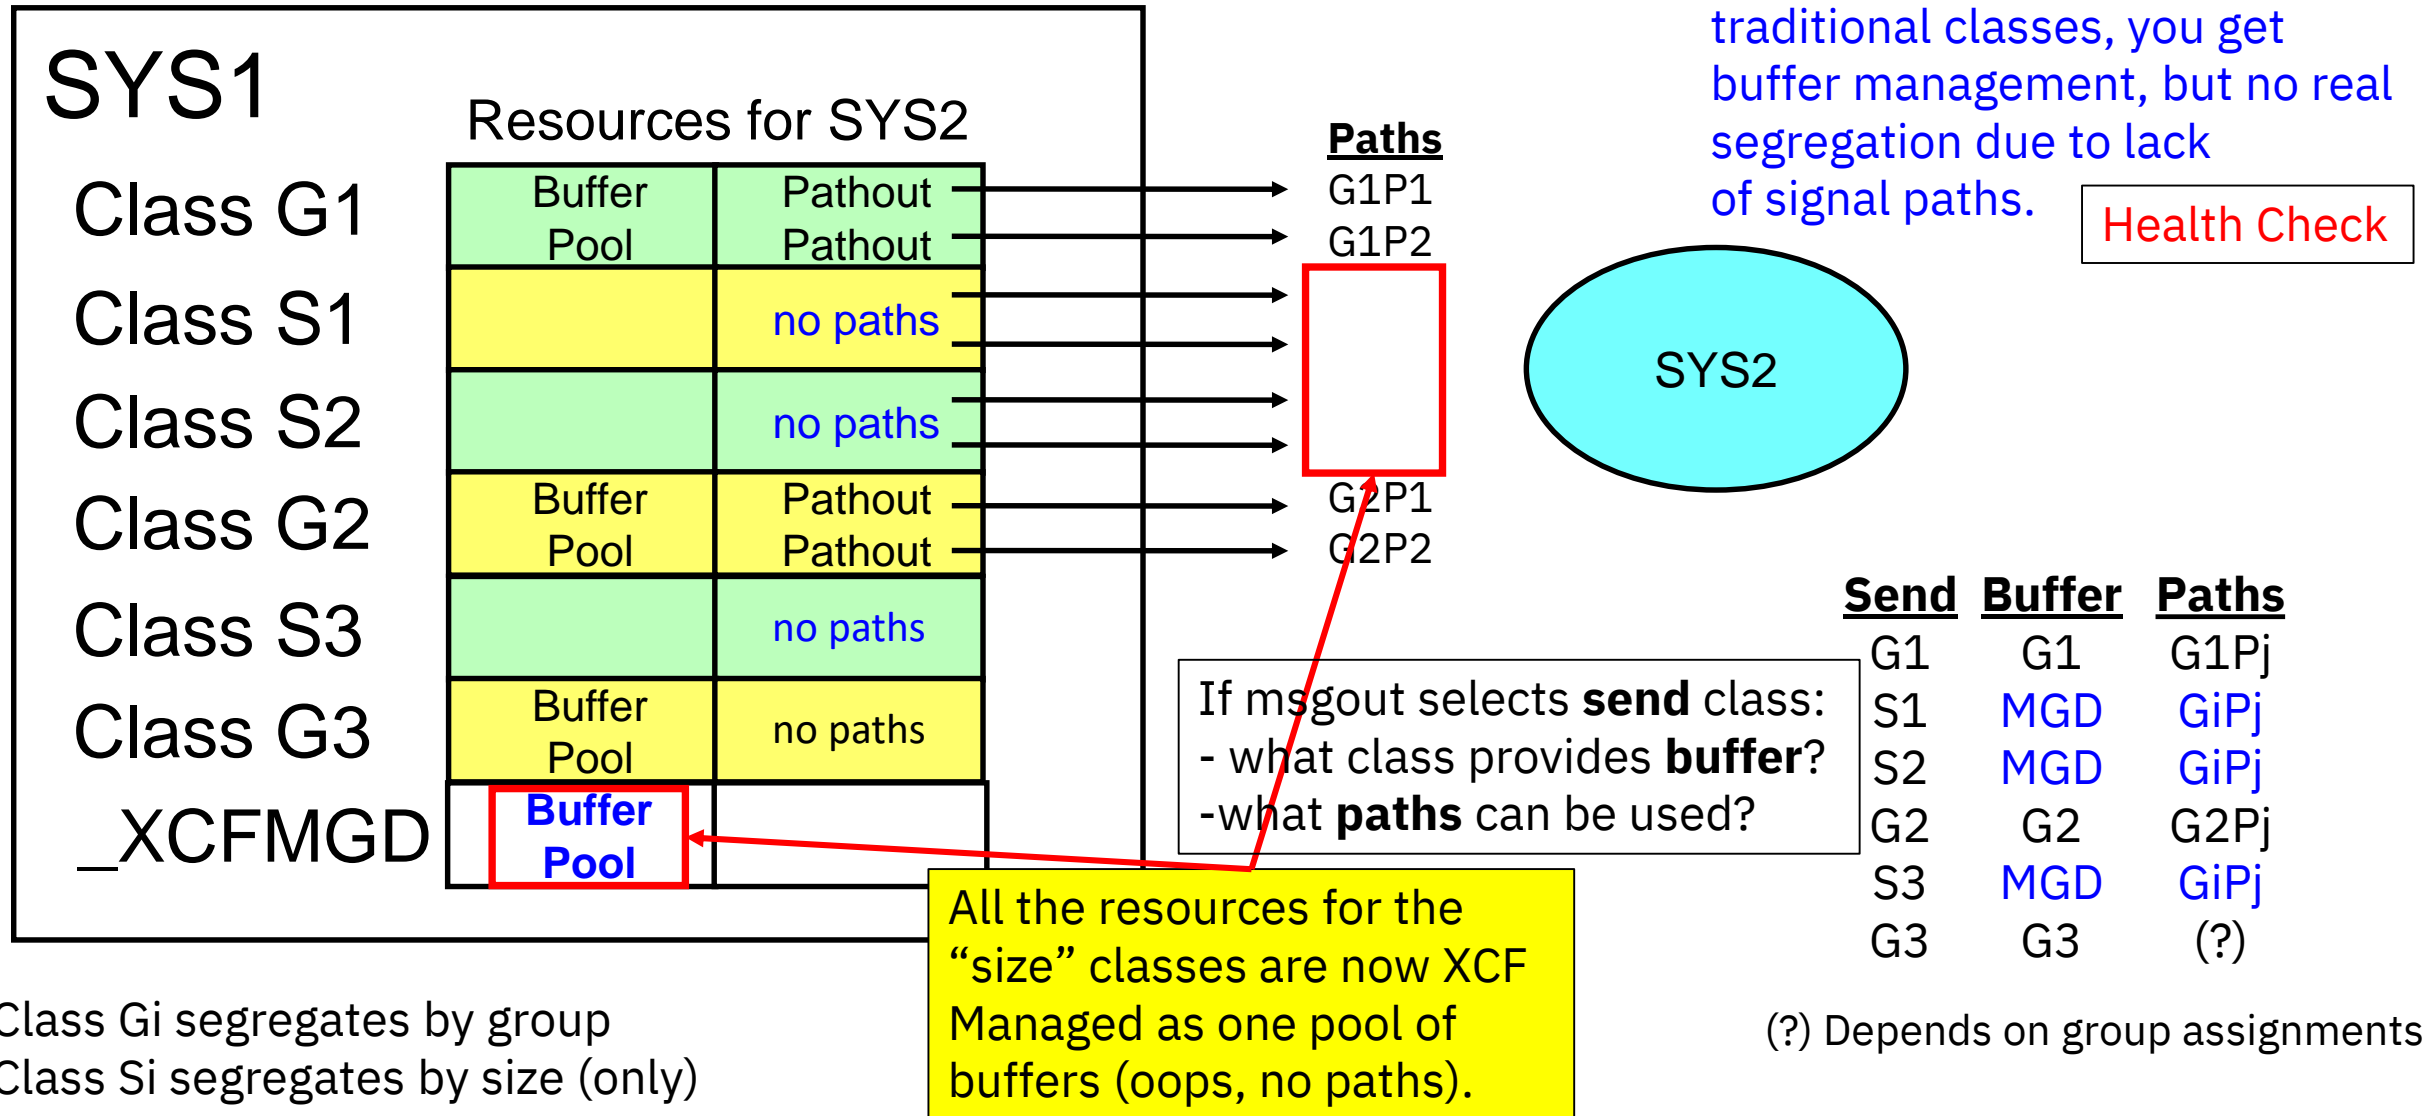
# Transport Classes when XTCSIZE is DISABLED



# Transport Classes when XTCSIZE is ENABLED



# Transport Classes when XTCSIZE is ENABLED



# About the \_XCFMGD pseudo transport class

- Uses “best fit” buffers on the send side
  - Maximizes number of signals that can be accepted for a given MAXMSG limit
    - Which is important for handling bursts and delays
  - Traditional classes generally use the defined size which might not be best fit
    - So could encounter “no buffer” condition sooner than with \_XCFMGD
- Paths run at the maximum signal size
  - So any size signal can be transmitted without any additional overhead
    - Never any need to re-negotiate (or tune) the paths
  - But that implies buffers on target system are likely bigger than needed
    - Which raises “no buffer” concerns (resiliency, capacity)

# Oversize inbound buffers?

- You can actually sustain high signal rates with very few buffers, provided those buffers are recycled quickly
  - For a very long time, we've been doing "buffer migration" to reclaim our buffers when member message delivery takes "too long" to free them
  - As of z/OS V2R2, "Message Isolation" allows us to do "aggressive migration" to reclaim more buffers more quickly
  - With z/OS V2R4, we now have "expeditious buffer replenishment" to eliminate much of the latency that elongated "no buffer" resolution
- So we believe we can now sustain high signal rates despite the potential for having fewer inbound buffers for given MAXMSG

# No significant increase in “no buffer” conditions

**DISABLED**

|             |                                    | INBOUND TO CB88 |                 |               |  |
|-------------|------------------------------------|-----------------|-----------------|---------------|--|
| FROM SYSTEM | T FROM/TO Y DEVICE, OR P STRUCTURE | REQ IN          | BUFFERS UNAVAIL | TRANSFER TIME |  |
| CB86        | S IXCPLEX_PATH1                    | 1060K           | 210             | 0.156         |  |
|             | S IXCPLEX_PATH2                    | 50              | 0               | 0.313         |  |
|             | S IXCPLEX_PATH3                    | 1337K           | 289             | 0.139         |  |
|             | S IXCPLEX_PATH4                    | 151             | 0               | 0.190         |  |
|             | C C580 TO C564                     | 62,268          | 0               | 0.196         |  |
|             | C C581 TO C565                     | 61,871          | 1,769           | 0.179         |  |
|             | C C582 TO C566                     | 624             | 0               | 0.227         |  |

2,521,964 REQ IN  
2,268 BUFFERS UNAVAIL

**ENABLED**

|      |                 |        |       |       |  |
|------|-----------------|--------|-------|-------|--|
| CB86 | S IXCPLEX_PATH1 | 109    | 0     | 0.310 |  |
|      | S IXCPLEX_PATH2 | 3,220  | 93    | 0.280 |  |
|      | S IXCPLEX_PATH3 | 1767K  | 2,333 | 0.131 |  |
|      | S IXCPLEX_PATH4 | 1,985  | 0     | 0.143 |  |
|      | C C580 TO C564  | 5,897  | 0     | 0.173 |  |
|      | C C581 TO C565  | 85,084 | 0     | 0.179 |  |
|      | C C582 TO C566  | 70,074 | 209   | 0.202 |  |

1,933,369 REQ IN  
2,635 BUFFERS UNAVAIL

Yes, “no buffer” increased.  
But certainly not to the extent one might expect based on relative buffer sizes.  
And not enough to impact throughput.



# Oversize inbound buffers - implications

- More buffer space consumed by the normal steady state working set (likely fewer in number, but they are larger)
  - Inbound list paths will drop to zero buffers in use if there is not enough traffic to keep them busy.
  - Inbound CTC devices generally have at least 4 buffers always in use.
- But may not need as much space for the peaks
  - We will still consume up to the MAXMSG limit you provide
  - For those that have been increasing inbound MAXMSG to avoid “no buffer” conditions, you may not need such large values
    - Indeed, large values can induce long queue effects
    - In some of our tests, throughput degraded as MAXMSG increased!

# Inbound “no buffer” conditions

- Our use of large buffers implies “fewer” buffers for a given MAXMSG value, which increases potential for “no buffer”
- Historically, “no buffer” seems to incent people to increase their MAXMSG values, sometimes to unreasonably high values
  - Generally, if you hit MAXMSG limit, you have queueing. The bigger the MAXMSG limit that you hit, the more queueing you have.
  - Driving “no buffer” to zero looks nice, but that queue containing tens of thousands of signals could be more detrimental to system performance than the “no buffer” condition itself
- On the inbound side, “no buffer” is not necessarily bad ....

# Inbound “no buffer” counts

- For a CTC device, I/O is started in anticipation of there being a signal to receive. If “no buffer”, we could not start the desired I/O.
  - But if there is nothing to receive, there is no signal transfer delay.
  - So long as we get a buffer and start the I/O before a signal is sent, there is no impact no matter how high the “no buffer” count climbs
- For a list path, I/O is started when we believe there is a signal to receive. So a “no buffer” condition likely implies delay.
  - But not necessarily. Sometimes our belief is unfounded.
- XCF abhors inbound “no buffer” conditions, so it repeatedly looks to resolve them, which can drive up the “no buffer” count.
- But the count tells us very little about the impact on signal transfer since we can’t tell how much delay was induced by the “no buffer” condition.

# Measuring impactful “no buffer” conditions

- We are providing additional inbound path measurements to capture the notion of an “impactful no buffer condition”
  - If we hit a “no buffer” condition, did it impact signal transfer?
  - If so, how long did the impact last?
- To some degree, we wanted insight as to whether our use of largest signal buffers (and therefore smaller number) was inducing signal transfer delays.
- But more generally, we anticipate this data will inform judgments about how many signal paths are needed:
  - Maybe we need more buffers instead of a new path
  - Maybe we need buffers to get recycled more quickly
  - Maybe we don’t care because the impacts are inconsequential

SMF data gathering  
and RMF reports are  
pending.

# XCF defines the \_XCFMGD pseudo transport class

- You cannot directly modify its attributes. However:
  - Buffer space MAXMSG limit used for \_XCFMGD takes into account:
    - Default MAXMSG value
    - Buffer space limits of the classes that are being managed
    - Our own judgment as to what is a reasonable value
  - So changes to default MAXMSG value, or MAXMSG values for a managed class or a managed path could cause the buffer limits for \_XCFMGD to change
    - Though you might need to DISABLE and then ENABLE the XTCSIZE switch to have them take effect
- You cannot directly assign paths to the \_XCFMGD class
  - Paths are “inherited” from the classes managed when XTCSIZE is ENABLED
  - So you continue to assign paths to a traditional transport class
  - But they are reported as \_XCFMGD when XTCSIZE is ENABLED
    - This lets you “see” the new behavior

DISPLAY XCF accepts \_XCFMGD as a class name.  
SETXCF and COUPLExx do not.

# What will you see?

DISPLAY XCF output

RMF reports of XCF activity

# D XCF,C – What is the XTCSIZE switch setting?

D XCF,C

IXC357I 16.08.07 DISPLAY XCF

SYSTEM SY1 DATA

....

## OPTIONAL FUNCTION STATUS:

| FUNCTION NAME   | STATUS         | DEFAULT        |
|-----------------|----------------|----------------|
| DUPLEXCF16      | DISABLED       | DISABLED       |
| SYSSTATDETECT   | ENABLED        | ENABLED        |
| USERINTERVAL    | ENABLED        | DISABLED       |
| CRITICALPAGING  | DISABLED       | DISABLED       |
| DUPLEXCFDIAG    | DISABLED       | DISABLED       |
| CFLCRMGMT       | DISABLED       | DISABLED       |
| COUPLINGTHININT | ENABLED        | ENABLED        |
| CFSTRQMON       | DISABLED       | DISABLED       |
| MSGISO          | ENABLED        | ENABLED        |
| <b>XTCSIZE</b>  | <b>ENABLED</b> | <b>ENABLED</b> |

....

XCF Managed is active on this system.

When sending signals to a target system running z/OS V2R4, the signal resources for all of the “size only” transport classes are being XCF Managed as a single pool that supports any size signal.

When sending signals to a down level target system, traditional transport classes and traditional segregation rules are used.

# D XCF,C – What is the XTCSIZE switch setting?

D XCF,C

IXC357I 16.08.07 DISPLAY XCF

SYSTEM SY1 DATA

....

## OPTIONAL FUNCTION STATUS:

| FUNCTION NAME   | STATUS          | DEFAULT        |
|-----------------|-----------------|----------------|
| DUPLEXCF16      | DISABLED        | DISABLED       |
| SYSSTATDETECT   | ENABLED         | ENABLED        |
| USERINTERVAL    | ENABLED         | DISABLED       |
| CRITICALPAGING  | DISABLED        | DISABLED       |
| DUPLEXCFDIAG    | DISABLED        | DISABLED       |
| CFLCRMGMT       | DISABLED        | DISABLED       |
| COUPLINGTHININT | ENABLED         | ENABLED        |
| CFSTRQMON       | DISABLED        | DISABLED       |
| MSGISO          | ENABLED         | ENABLED        |
| <b>XTCSIZE</b>  | <b>DISABLED</b> | <b>ENABLED</b> |

....

XCF Managed is not active.

You are using traditional transport classes and traditional transport class segregation rules.



# Change XTCSIZE switch setting - dynamically

**SETXCF FUNCTIONS,DISABLE=XTCSIZE**

Revert to old behavior

IXC373I XCF / XES OPTIONAL FUNCTIONS DISABLED:  
XTCSIZE

**SETXCF FUNCTIONS,ENABLE=XTCSIZE**

Enable new behavior

IXC373I XCF / XES OPTIONAL FUNCTIONS ENABLED:  
XTCSIZE

# Change XTCSIZE switch setting – through IPL

## COUPLExx parmlib member

COUPLE    SYSPLEX(plexname)

PCOUPLE(...)

ACOUPLE(...)

...

## **FUNCTIONS ENABLE(XTCSIZE)**

CLASSDEF ...            Enable new behavior

PATHIN    ...

PATHOUT    ...

## COUPLExx parmlib member

COUPLE    SYSPLEX(plexname)

PCOUPLE(...)

ACOUPLE(...)

...

## **FUNCTIONS DISABLE(XTCSIZE)**

CLASSDEF ...            Revert to old behavior

PATHIN    ...

PATHOUT    ...

# D XCF,CLASSDEF – to see transport class data

```
D XCF,CLASSDEF,CLASS=ALL
```

```
IXC344I 16.08.23 DISPLAY XCF 848
```

| TRANSPORT      | CLASS    | DEFAULT     | ASSIGNED       |
|----------------|----------|-------------|----------------|
| CLASS          | LENGTH   | MAXMSG      | GROUPS         |
| <b>_XCFMGD</b> | <b>0</b> | <b>2000</b> | <b>UNDESIG</b> |
| BADGUY0        | 956      | 2000        | MYAPP1         |
| BADGUY1        | 4028     | 2000        | MYAPP1 UNDESIG |
| BADGUY2        | 8124     | 2000        | UNDESIG        |
| BIG            | 32000    | 2000        | UNDESIG        |
| DEFAULT        | 956      | 2000        | UNDESIG        |
| MED            | 12000    | 2000        | UNDESIG        |
| SML            | 956      | 2000        | UNDESIG        |

Always  
present →

These classes  
selected when  
XTCSIZE is  
ENABLED.

We don't have anything to show you which classes qualify for XTCSIZE management.

# D XCF,CLASSDEF – to see transport class data (cont)

As for any traditional class, \_XCFMGD signal size distributions are shown for each target system.

...

\_XCFMGD TRANSPORT CLASS USAGE FOR SYSTEM SY1

|             |       |                |      |         |   |
|-------------|-------|----------------|------|---------|---|
| SUM MAXMSG: | 10000 | IN USE:        | 2    | NOBUFF: | 0 |
| SEND CNT:   | 978   | BUFFLEN (FIT): | 956  |         |   |
| SEND CNT:   | 362   | BUFFLEN (BIG): | 4028 |         |   |
| SEND CNT:   | 21    | BUFFLEN (BIG): | 8124 |         |   |

SY1 is local system

Should we say FIT?

\_XCFMGD TRANSPORT CLASS USAGE FOR SYSTEM SY2

|             |       |                |      |         |   |
|-------------|-------|----------------|------|---------|---|
| SUM MAXMSG: | 22000 | IN USE:        | 8    | NOBUFF: | 0 |
| SEND CNT:   | 1713  | BUFFLEN (FIT): | 956  |         |   |
| SEND CNT:   | 130   | BUFFLEN (BIG): | 4028 |         |   |
| SEND CNT:   | 119   | BUFFLEN (BIG): | 8124 |         |   |

SY2 running z/OS V2R4  
("XCF Managed" is in play)

\_XCFMGD TRANSPORT CLASS USAGE FOR SYSTEM SY3

|             |       |                |     |         |   |
|-------------|-------|----------------|-----|---------|---|
| SUM MAXMSG: | 16000 | IN USE:        | 0   | NOBUFF: | 0 |
| SEND CNT:   | 0     | BUFFLEN (FIT): | 956 |         |   |

SY3 running z/OS V2R3  
Down level, so old rules.  
("XCF Managed" not used)

# Which list paths are XCF Managed ?

D XCF,P0,STRNAME=ALL,STATUS=WORKING

IXC356I 16.09.28 DISPLAY XCF 855

| STRNAME        | REMOTE | PATHOUT | UNUSED | TRANSPORT |        |                |
|----------------|--------|---------|--------|-----------|--------|----------------|
| PATHOUT        | SYSTEM | STATUS  | PATHS  | RETRY     | MAXMSG | CLASS          |
| IXCTL_SIGNAL01 |        | WORKING | 234    | 10        | 2000   | DEFAULT        |
|                | SY2    | WORKING |        |           |        | <u>_XCFMGD</u> |
|                | SY3    | WORKING |        |           |        | DEFAULT        |
| IXCTL_SIGNAL02 |        | WORKING | 234    | 10        | 2000   | DEFAULT        |
|                | SY2    | WORKING |        |           |        | _XCFMGD        |
|                | SY3    | WORKING |        |           |        | DEFAULT        |

Structure still indicates the "home" class

Uplevel target system

Downlevel target system (so using old rules)

| STRNAME        | REMOTE | PATHOUT | TRANSFR | BUFFER  | MSGBUF       | SIGNL  | MXFER |      |
|----------------|--------|---------|---------|---------|--------------|--------|-------|------|
| PATHOUT        | LIST   | SYSTEM  | STATUS  | PENDING | LENGTH       | IN USE | NUMBR | TIME |
| IXCTL_SIGNAL01 |        |         |         |         |              |        |       |      |
|                | 12     | SY2     | WORKING | 0       | <u>62464</u> | 204    | 10    | 2915 |
|                | 9      | SY3     | WORKING | 0       | 956          | 6      | 3     | 4612 |
| IXCTL_SIGNAL02 |        |         |         |         |              |        |       |      |
|                | 12     | SY2     | WORKING | 0       | 62464        | 270    | 11    | 2767 |
|                | 9      | SY3     | WORKING | 0       | 956          | 8      | 4     | 1049 |

Running at biggest size

Running at defined size

# Same list paths after XTCSIZE is DISABLED

D XCF,P0,STRNAME=ALL,STATUS=WORKING

No class is XCF Managed

IXC356I 16.10.32 DISPLAY XCF 871

| STRNAME        | REMOTE | PATHOUT | UNUSED | TRANSPORT |        |         |
|----------------|--------|---------|--------|-----------|--------|---------|
| PATHOUT        | SYSTEM | STATUS  | PATHS  | RETRY     | MAXMSG | CLASS   |
| IXCTL_SIGNAL01 |        | WORKING | 234    | 10        | 2000   | DEFAULT |
|                | SY2    | WORKING |        |           |        | DEFAULT |
|                | SY3    | WORKING |        |           |        | DEFAULT |
| IXCTL_SIGNAL02 |        | WORKING | 234    | 10        | 2000   | DEFAULT |
|                | SY2    | WORKING |        |           |        | DEFAULT |
|                | SY3    | WORKING |        |           |        | DEFAULT |

Reverts to "home" class

| STRNAME        | REMOTE | PATHOUT | TRANSFR | BUFFER  | MSGBUF | SIGNL  | MXFER |      |
|----------------|--------|---------|---------|---------|--------|--------|-------|------|
| PATHOUT        | LIST   | SYSTEM  | STATUS  | PENDING | LENGTH | IN USE | NUMBR | TIME |
| IXCTL_SIGNAL01 |        |         |         |         |        |        |       |      |
|                | 12     | SY2     | WORKING | 0       | 956    | 12     | 18    | 3418 |
|                | 9      | SY3     | WORKING | 0       | 956    | 8      | 4     | 3044 |
| IXCTL_SIGNAL02 |        |         |         |         |        |        |       |      |
|                | 12     | SY2     | WORKING | 0       | 956    | 14     | 19    | 262  |
|                | 9      | SY3     | WORKING | 0       | 956    | 10     | 5     | 3116 |

Reverts to defined size

# Which CTC signal paths are XCF Managed?

DEFAULT class is XCF Managed.  
BADGUY0 is not.

```
D XCF,PO,DEV=ALL,STATUS=WORKING
IXC356I 16.09.14 DISPLAY XCF 852
```

| LOCAL DEVICE | REMOTE     | PATHOUT        | REMOTE      | RETRY     | MAXMSG      | TRANSPORT      |
|--------------|------------|----------------|-------------|-----------|-------------|----------------|
| PATHOUT      | SYSTEM     | STATUS         | PATHIN      |           |             | CLASS          |
| 8000         | SY2        | WORKING        | 80BF        | 10        | 2000        | BADGUY0        |
| <b>8001</b>  | <b>SY2</b> | <b>WORKING</b> | <b>80BE</b> | <b>10</b> | <b>2000</b> | <b>_XCFMGD</b> |
| 8002         | SY3        | WORKING        | 80BD        | 10        | 2000        | BADGUY0        |
| 8003         | SY3        | WORKING        | 80BC        | 10        | 2000        | DEFAULT        |

You can't see the  
"home" class

Uplevel target system

Downlevel target system  
(so using old rules)

| LOCAL       | REMOTE      | REMOTE     | PATHOUT        | TRANSFR  | BUFFER       | MSGBUF    | SIGNL       | MXFER     |
|-------------|-------------|------------|----------------|----------|--------------|-----------|-------------|-----------|
| PATHOUT     | PATHIN      | SYSTEM     | STATUS         | PENDING  | LENGTH       | IN USE    | NUMBR       | TIME      |
| 8000        | 80BF        | SY2        | WORKING        | 0        | 956          | 14        | 826         | 34        |
| <b>8001</b> | <b>80BE</b> | <b>SY2</b> | <b>WORKING</b> | <b>0</b> | <b>62464</b> | <b>14</b> | <b>1387</b> | <b>38</b> |
| 8002        | 80BD        | SY3        | WORKING        | 0        | 956          | 14        | 1193        | 48        |
| 8003        | 80BC        | SY3        | WORKING        | 0        | 956          | 14        | 1847        | 29        |

Running at biggest size

Running at defined size

# Same CTC paths after XTCSIZE is DISABLED

No class is XCF Managed

```
D XCF,PO,DEV=ALL,STATUS=WORKING
IXC356I 16.10.28 DISPLAY XCF 868
```

| LOCAL DEVICE | REMOTE | PATHOUT | REMOTE | RETRY | MAXMSG | TRANSPORT |
|--------------|--------|---------|--------|-------|--------|-----------|
| PATHOUT      | SYSTEM | STATUS  | PATHIN |       |        | CLASS     |
| 8000         | SY2    | WORKING | 80BF   | 10    | 2000   | BADGUY0   |
| 8001         | SY2    | WORKING | 80BE   | 10    | 2000   | DEFAULT   |
| 8002         | SY3    | WORKING | 80BD   | 10    | 2000   | BADGUY0   |
| 8003         | SY3    | WORKING | 80BC   | 10    | 2000   | DEFAULT   |

Reverts to "home" class

| LOCAL   | REMOTE | REMOTE | PATHOUT | TRANSFR | BUFFER | MSGBUF | SIGNL | MXFER |
|---------|--------|--------|---------|---------|--------|--------|-------|-------|
| PATHOUT | PATHIN | SYSTEM | STATUS  | PENDING | LENGTH | IN USE | NUMBR | TIME  |
| 8000    | 80BF   | SY2    | WORKING | 0       | 956    | 14     | 828   | 35    |
| 8001    | 80BE   | SY2    | WORKING | 0       | 956    | 14     | 1497  | 36    |
| 8002    | 80BD   | SY3    | WORKING | 0       | 956    | 14     | 1194  | 29    |
| 8003    | 80BC   | SY3    | WORKING | 0       | 956    | 14     | 1950  | 23    |

Reverts to defined size



# RMF Report: XCF Usage by System

| OUTBOUND FROM S5A   |                    |                  |            |                    |          |          |          |                         |               |
|---------------------|--------------------|------------------|------------|--------------------|----------|----------|----------|-------------------------|---------------|
| TO<br>SYSTEM<br>S5B | TRANSPORT<br>CLASS | BUFFER<br>LENGTH | REQ<br>OUT | ----- BUFFER ----- |          |          |          | ALL<br>PATHS<br>UNAVAIL | REQ<br>REJECT |
|                     |                    |                  |            | %<br>SML           | %<br>FIT | %<br>BIG | %<br>OVR |                         |               |
|                     | _XCFMGD            | 956              | 624,950    | 0                  | 99       | 1        | 100      | 0                       | 0             |
|                     | BIG                | 40,892           | 0          | 0                  | 0        | 0        | 0        | 0                       | 0             |
|                     | CTTX               | 40,892           | 0          | 0                  | 0        | 0        | 0        | 0                       | 0             |
|                     | DAE                | 956              | 0          | 0                  | 0        | 0        | 0        | 0                       | 0             |
|                     | DEFAULT            | 20,412           | 0          | 0                  | 0        | 0        | 0        | 0                       | 0             |
|                     | DEFSMALL           | 956              | 0          | 0                  | 0        | 0        | 0        | 0                       | 0             |
|                     | DEF8K              | 8,124            | 0          | 0                  | 0        | 0        | 0        | 0                       | 0             |
|                     | FEWFAST            | 956              | 0          | 0                  | 0        | 0        | 0        | 0                       | 0             |

\_XCFMGD transport class will always appear.

- Its counts will be zero if XTCSIZE is DISABLED or if the target system is down level.
- SML, FIT, BIG are truthful and consistent with existing definitions of SMF data.

In the past, these numbers might have prompted one to consider revising class definitions.

**But with \_XCFMGD, not only is there nothing to be done, there is nothing you can do.**

The traditional classes will always appear as well. The counts will be zero if the class is not being used. But from this report, you can't tell why there's no use: XCF Managed? No signals sent? (CTTX for example).

# RMF Report: XCF Path Statistics

| OUTBOUND FROM S5A |        |                                    |                    |            |               |         |      |       |  |
|-------------------|--------|------------------------------------|--------------------|------------|---------------|---------|------|-------|--|
| TO<br>SYSTEM      | T<br>Y | FROM/TO<br>DEVICE, OR<br>STRUCTURE | TRANSPORT<br>CLASS | REQ<br>OUT | AVG Q<br>LNTH | AVAIL   | BUSY | RETRY |  |
| S5B               | S      | IXCPLEX_PATH1                      | _XCFMGD            | 147,251    | 0.00          | 147,251 | 0    | 0     |  |
|                   | S      | IXCPLEX_PATH2                      | _XCFMGD            | 454,322    | 0.00          | 454,322 | 0    | 0     |  |
|                   | S      | IXCPLEX_PATH3                      | _XCFMGD            | 5,011      | 0.00          | 5,011   | 0    | 0     |  |
|                   | S      | IXCPLEX_PATH4                      | CTTX               | 0          | 0.00          | 0       | 0    | 0     |  |
|                   | S      | IXCPLEX_PATH5                      | _XCFMGD            | 18,395     | 0.00          | 18,383  | 12   | 0     |  |
|                   | C      | C5B0 TO C5A4                       | _XCFMGD            | 57         | 0.00          | 53      | 4    | 0     |  |
|                   | C      | C5B1 TO C5A5                       | _XCFMGD            | 73         | 0.00          | 64      | 9    | 0     |  |
|                   | C      | C5B2 TO C5A6                       | _XCFMGD            | 52         | 0.00          | 47      | 5    | 0     |  |



Paths are reported as being assigned to \_XCFMGD transport class if they are being XCF Managed.

# Our exit criteria

## **Assuming XTCSIZE is ENABLED and target system is z/OS V2R4:**

- The (pure size) Transport Class definitions will be completely irrelevant to performance of the sysplex workload
  - Regardless of what you specify for Transport Class definitions and regardless of how you assign signal paths to them {and specify MAXMSG values?}, there will not be any noticeable impact on the sysplex workload
  - No matter how “good” or “bad” those specifications might be
  - Changes to class definitions will neither improve nor degrade signal delivery
- The only determining factors for signal performance will be:
  - The number of signal paths you provide
  - The performance characteristics of those paths
  - The performance characteristics of the systems using those paths

# Really not too tough a comparison to win

- Most of you likely have:
  - Multiple “pure size” transport classes defined
  - One or more signal paths assigned to each of those classes
  - Very high AVAIL percentages
    - Signal paths nearly always “not busy” when picked to send a signal
    - Some paths hardly ever picked, so lots of excess capacity
  - Hardly any “no buffer” conditions
- When XTCSIZE is ENABLED
  - Best fit buffers for \_XCFMGD likely implies more send side buffer capacity than your definitions provide
  - Signals likely distributed to paths much as your traditional classes do today
  - An underutilized path in one of your classes can now be used to help a peer class that needs more capacity

# My worry

- Maybe your classes look like they are “pure size”
- But there is some “ill behaved” group that only sends one size msg
  - So those messages land in one of your classes
  - That is seldom if ever used by other groups
- That is, this “size only” class is effectively providing group segregation for the ill behaved group
- When XTCSIZE is ENABLED, that ill behaved group is now released into the general population where it can infect the world. That could be bad.
  - Existing “msg isolation” might save the day
  - If you recognize the problem, you could define a class to do the needed group segregation
  - While we wait for XCF to eliminate the need to segregate signals by group.

# Our expectations

- Overall, your sysplex workloads should run at least as well as they do now. Signal performance should not be degraded.
  - Regardless of what z/OS release the target system is running
  - Regardless of whether XTCSIZE is ENABLED or DISABLED
- For a given set of signal paths, there is no set of (pure size) transport class definitions for those paths that will perform better with XTCSIZE=DISABLED than with XTCSIZE=ENABLED.
  - In other words, if you have signaling related problem that can be resolved with traditional (pure size) transport classes using the same (or even fewer) number of signal paths, we failed.

# Caveats

**I really don't expect there to be much change to the behavior of your workloads. But there are possibilities:**

- Our overall improvements to signal delivery might change timings that induce secondary impacts
- Signal traffic might flow over different paths than in the past, which might impact things like CF, subchannel, and link utilization
  - If you have classes with “busy” paths today, those signals are likely to be sent via an available path from a peer XCF Managed class.
  - So maybe your high frequency signals always traveled via CF1. With XTCSIZE=ENABLED, any signal can go via any XCF Managed path. So CF2, which in the past got very little activity, might now be used more frequently.

# Signal distributions might change

•

OUTBOUND FROM CB86

| TO SYSTEM       | T FROM/TO Y DEVICE, OR P STRUCTURE | TRANSPORT CLASS | REQ OUT | AVG Q LGTH | AVAIL   | BUSY   | RETRY |
|-----------------|------------------------------------|-----------------|---------|------------|---------|--------|-------|
| <b>DISABLED</b> | CB8C S IXCplex_PATH1               | DEFsmall        | 1150910 | 0.00       | 1084034 | 66,876 | 0     |
|                 | S IXCplex_PATH2                    | DEFMED          | 931     | 0.00       | 931     | 0      | 0     |
|                 | S IXCplex_PATH3                    | DEFsmall        | 3301531 | 0.00       | 3225472 | 76,059 | 0     |
|                 | S IXCplex_PATH4                    | DEFMED          | 30      | 0.00       | 30      | 0      | 0     |
| <b>ENABLED</b>  | CB8C S IXCplex_PATH1               | _XCfmGD         | 925,035 | 0.00       | 924,888 | 147    | 0     |
|                 | S IXCplex_PATH2                    | _XCfmGD         | 362,281 | 0.00       | 362,146 | 135    | 0     |
|                 | S IXCplex_PATH3                    | _XCfmGD         | 3362504 | 0.00       | 3362334 | 170    | 0     |
|                 | S IXCplex_PATH4                    | _XCfmGD         | 13,133  | 0.00       | 12,901  | 232    | 0     |

## With XTCSIZE ENABLED

- Signals more evenly distributed across paths
- Number of “busy” becomes negligible.
- (not shown) Inbound went from 0 to 6 no-buffer conditions.



# Session Summary

- We have eliminated the need for you to define XCF Transport Classes to segregate signals by size
  - No more planning, defining, monitoring, tuning, changing, ...
- Just **ENABLE** the **XTCSIZE** switch. We'll do the rest.
  - Which for ESP, is done for you. Sit back and enjoy the ride.
  - In particular, you don't need to change your current XCF configuration
  - One possible worry is potential impact on capacity if signals are distributed to paths in radically different patterns than what you see today
- We appreciate your feedback and the effort to help us compare **DISABLED** vs **ENABLED** in your environment.

# Appendix

- Contacts
  - Mark A Brooks – [mabrook@us.ibm.com](mailto:mabrook@us.ibm.com)
  - Neil Johnson - [najohnsn@us.ibm.com](mailto:najohnsn@us.ibm.com)