

z/OS 2.4 IBM Education Assistant (IEA)

Solution (Epic) Name: XCF Transport Class Simplification
Element(s)/Component(s): z/OS XCF



Trademarks

- See url <http://www.ibm.com/legal/copytrade.shtml> for a list of trademarks.

Session Objectives

- A little background on XCF Transport Classes to provide some context
- Where are we headed?
- What do you get in z/OS V2R4?
 - How enabled?
 - What changes?
 - Caveats and concerns?

Background

Brief review of Transport Class basics

Background: Transport Classes

- Transport classes segregate message traffic by size, group, or both
- Message traffic is segregated by size for the purpose of:
 - Ensuring timely transfer for small messages
 - Efficient utilization of buffer space
- Message traffic is segregated by group for the purpose of:
 - Isolating an ill-behaved group to protect others from harm
 - Dedicating XCF signal resources to a favored group so that it need not compete with others. Ostensibly, the group will:
 - Never be denied service
 - Never experience transfer delays

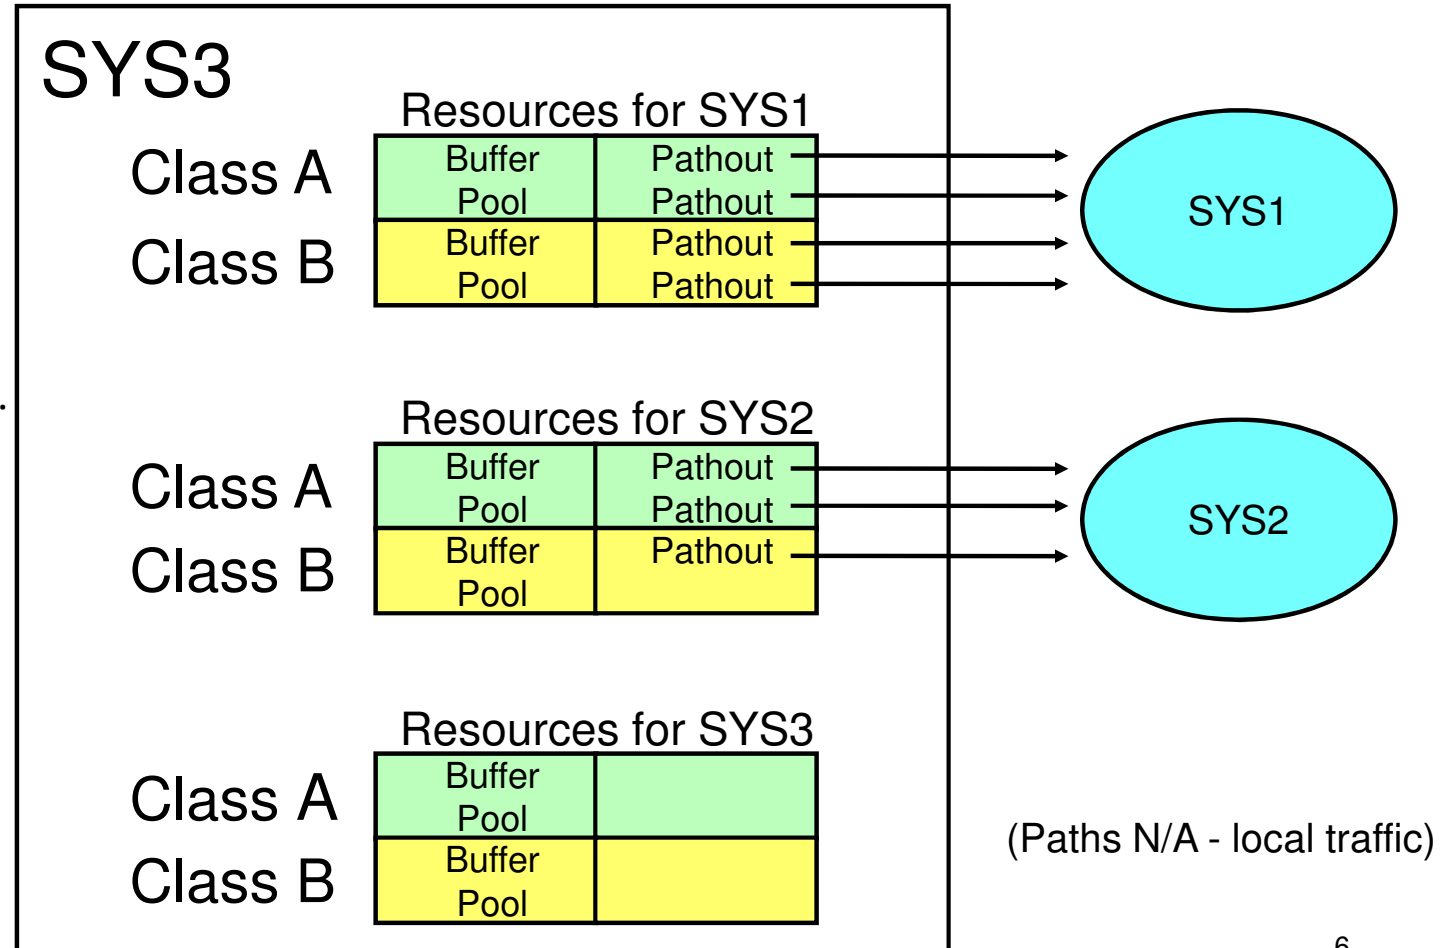
Background: Transport Classes

Transport class has buffers and signal paths (unless local).

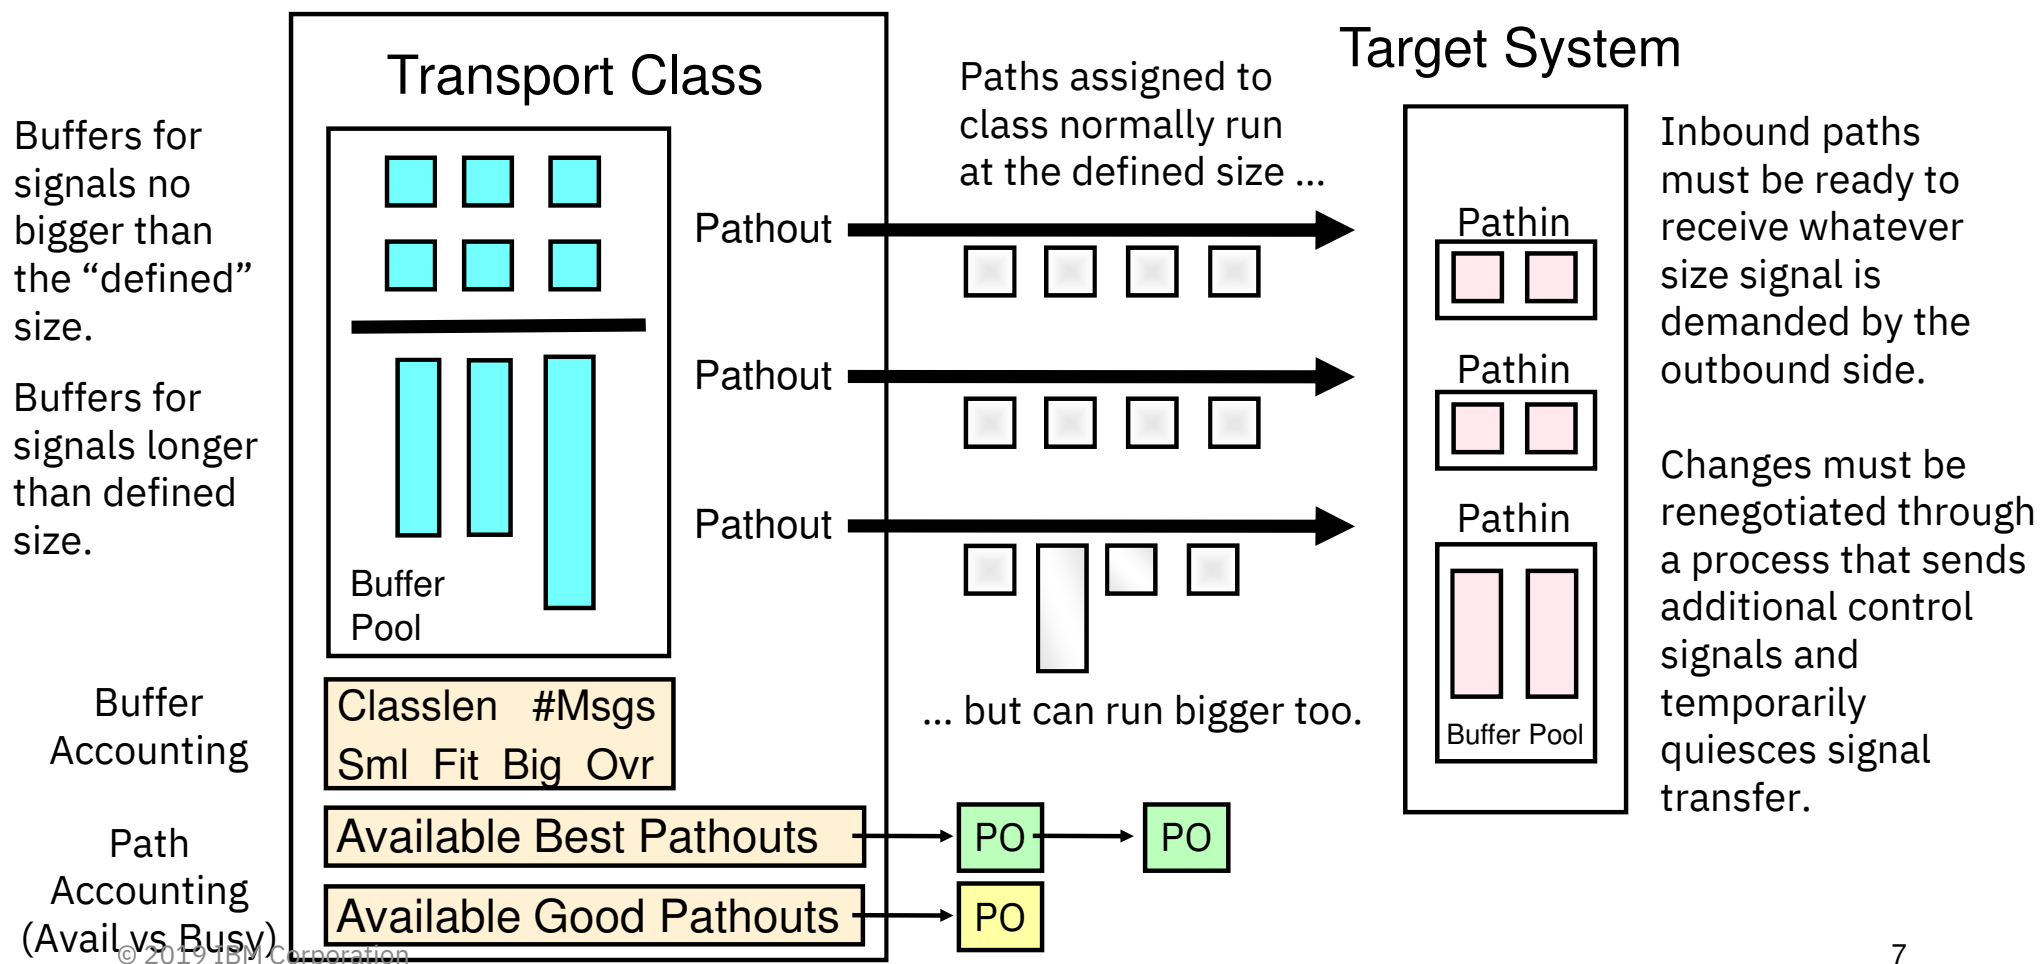
Each target system has it's own dedicated signal resources. Never shared between systems.

The same class definition applies to all target systems.

In practice, the signal traffic patterns are often vary dramatically on a target system basis.

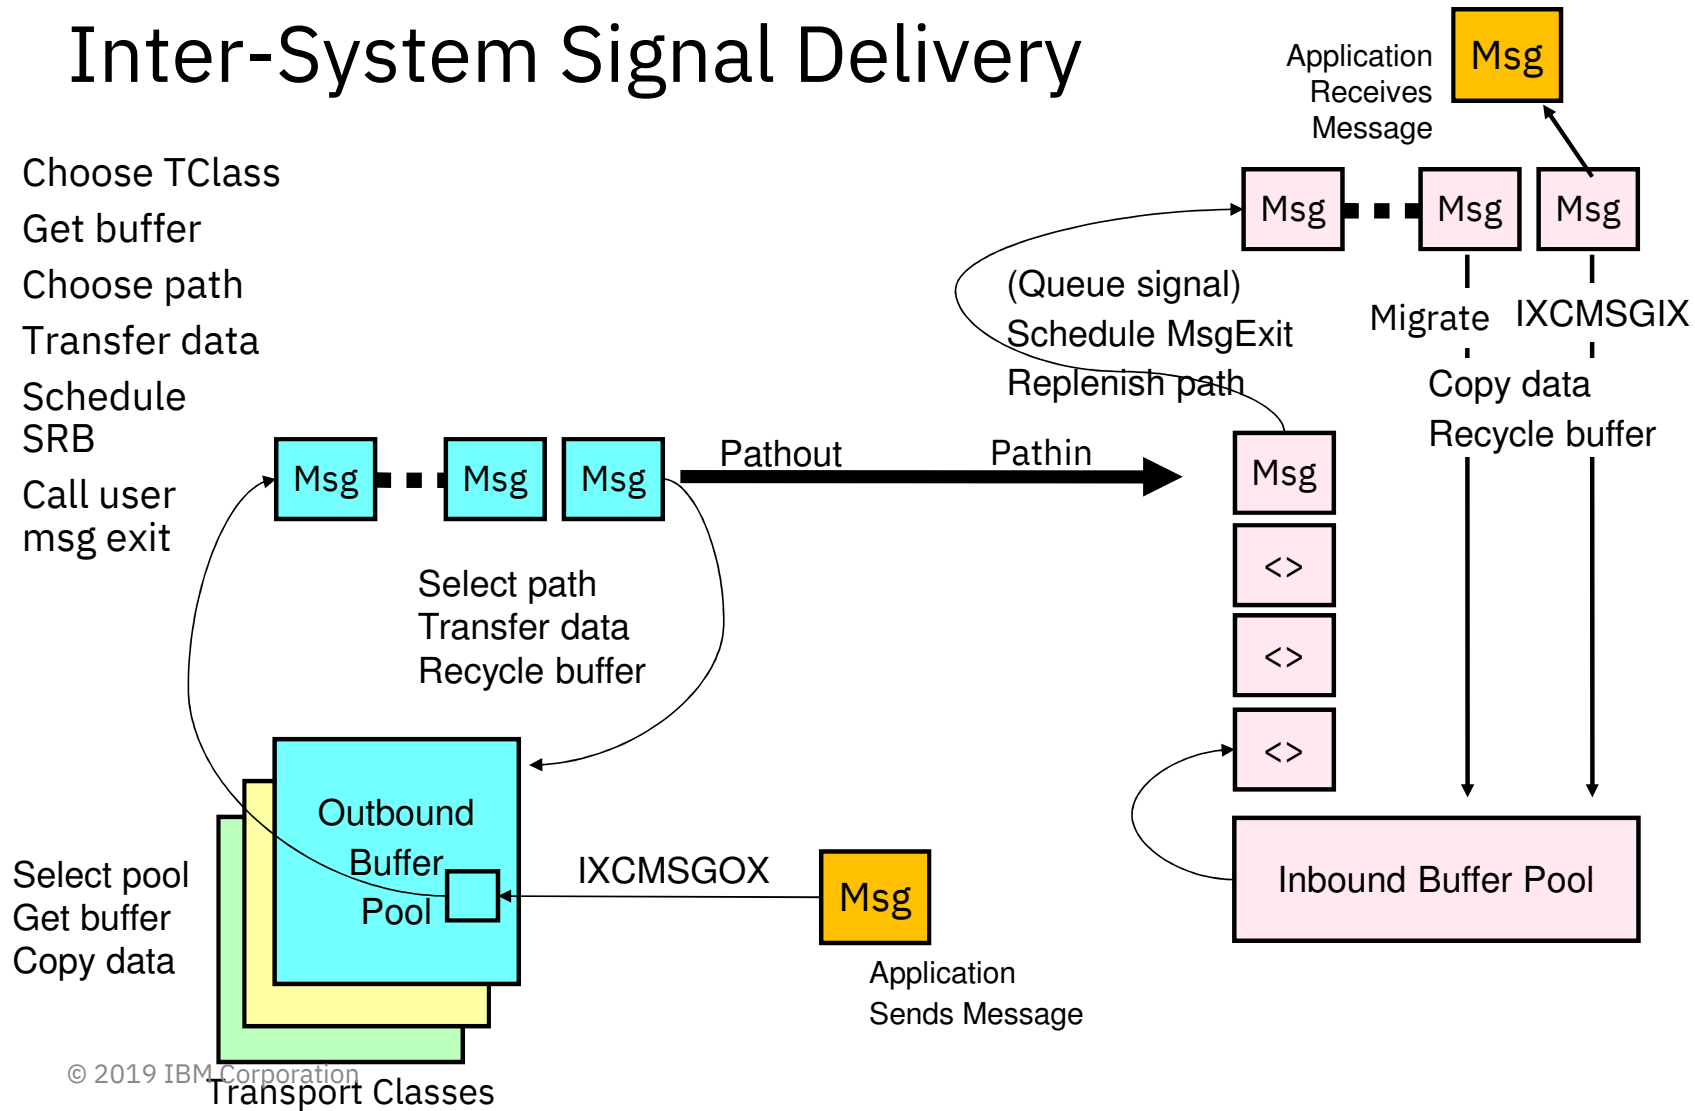


Background: Transport Class



Inter-System Signal Delivery

- Choose TClass
- Get buffer
- Choose path
- Transfer data
- Schedule SRB
- Call user msg exit



Where are we headed?

The installation can achieve the high level of resiliency, availability, and performance expected of sysplex without having to plan, define, monitor, or tune z/OS XCF transport classes.

Why eliminate the need to define Transport Classes?

- Simplification
 - Remove this burden from customers and all the IBM personnel that have to analyze, explain, make recommendations, etc.
 - Hard to provide clear guidance as to how classes should be configured; more art than science
- Avoid Outages
 - Transport Classes often not well understood, mistakes are made that lessen the resiliency of the sysplex and so permit avoidable outages
 - A Transport Class can help isolate an ill behaved member, but not really practical since you have to know in advance who is going to cause trouble
 - Such as delays and overwhelming bursts
- Self-optimization
 - Static definitions are not well suited to the dynamics of the sysplex
 - Can lead to inefficient use of resources (idle paths, excess storage)
 - System should automatically apply resources where needed most
 - For example, to handle a burst of messages or a change in signal patterns

To eliminate the need to define Transport Classes

XCF must automatically handle the problems that transport classes were intended to address:

- Timely message transfer
 - Maintain signal throughput and minimize signal delivery times, especially for the small messages that are typically most predominant in the sysplex
 - Provides better resiliency (minimize delays and queueing)
 - Efficient utilization of signal resources
 - Helps minimize cost
 - Provides better resiliency (more capacity)
-
- Isolation of ill behaved members
 - Avoid sympathy sickness so that problems with signal delivery for one member don't negatively impact signal delivery for other members
 - Fair access to signal resources
 - Don't allow one member to monopolize the signal resources to the detriment of signal delivery for other members

Size segregation

Group segregation

To eliminate the need to define Transport Classes

XCF must automatically handle the problems that transport classes were intended to address:

- Timely message transfer
 - Maintain signal throughput and minimize signal delivery times, especially for the small messages that are typically most predominant in the sysplex
 - Provides better resiliency

In the early days of sysplex, there was a significant difference in transfer time for small vs large signals. Today, this difference is much, much smaller and signal throughput is much greater. The relative delay arising from a small signal following a large one vs that of following a small signal is negligible. Intermixing signal sizes on a signal path will not have any discernable impact on the performance of the workload.

For z/OS V2R4, timely message transfer has been a primary focus. Our rationale:

- There will likely be some consolidation of signal paths as transport classes are eliminated, we want to make sure the survivors can sustain the signal load.
- What if there is a workload for which those marginal differences in transfer times for small vs large signals does matter? The ability to sustain higher signal rates will likely mask any impacts that might arise.
- And

To eliminate the need to define Transport Classes

XCF must automatically handle the problems that transport classes were intended to address:

- Timely message transfer
 - Maintain signal throughput and minimize signal delay for small messages that are typically most prevalent
 - Provides better resiliency
- Efficient utilization of signal resources
 - Helps minimize cost
 - Provides better resiliency
- Isolation of ill behaved members
 - Avoid sympathy sickness so that problems don't negatively impact signal delivery for other members
- Fair access to signal resources
 - Don't allow one member to monopolize the signal delivery for other members

These days, the more compelling reason for segregating signals by size has been to ensure that we can **maximize the number of signal buffers** available to the inbound path within the constraint of its MAXMSG limit. In general, more buffers tends to improve resiliency and helps maintain throughput.

Since our solution for eliminating the need to segregate signals by size uses the maximum size inbound buffer, we effectively **minimize the number of signal buffers** available to the inbound path. Even more impetus to focus on timely transfer and signal throughput!

To eliminate the need to define Transport Classes

XCF must automatically handle the problems that transport classes were intended to address:

- Timely message transfer
 - Maintain signal throughput and minimize signal delivery times, especially for the small messages that are typically most prevalent
 - Provides better resiliency
- Efficient utilization of signal resources
 - Helps minimize cost
 - Provides better resiliency
- Isolation of ill behaved members
 - Avoid sympathy sickness so that problem members don't negatively impact signal delivery for other members
- Fair access to signal resources
 - Don't allow one member to monopolize the signal delivery for other members

“Message Isolation” was delivered in z/OS V2R2. It is a foundational, incremental step along the journey towards the goal of eliminating the need to define transport classes for the purpose of segregating signals by group. Might provide “good enough” isolation, but it is not sufficient to make the claim that there is no need to define transport classes to segregate signals by group.

To eliminate the need to define Transport Classes

XCF must automatically handle the problems that transport classes were intended to address:

- Timely message transfer
 - Maintain signal throughput and minimize signal delivery times, especially for the small messages that are typically most problematic
 - Provides better resiliency
- Efficient utilization of signal resources
 - Helps minimize cost
 - Provides better resiliency
- Isolation of ill behaved members
 - Avoid sympathy sickness so that problem members don't negatively impact signal delivery for other members
- Fair access to signal resources
 - Don't allow one member to monopolize the signal delivery for other members

In general, we don't see field problems that would have been avoided if only we had had fair access to the XCF signal service. What I think of as "isolation" issues are more prevalent and problematic. But until we can guarantee fairness, we can't claim there is no need to define transport classes to segregate signals by group.

So what do you get?

Elimination of the need for transport classes to segregate signals by size.

(Elimination of need for group segregation is not included with this solution).

Overview

- Who (Audience)
 - System programmers, sysplex architects, performance analysts, capacity analysts, diagnosticians
- What (Solution)
 - XCF will internally manage the signal resources to provide timely delivery of signals in a sysplex independent of any Transport Class definitions **created purely for size segregation**.
- Wow (Benefit / Value, Need Addressed)
 - Simplification: You no longer need to define, monitor, tune, or manage XCF Transport Class definitions to segregate signals **purely by size**.
 - Simplification: You need only configure an appropriate number of signal paths.
 - Resilience: Less potential for non-optimal transport class definitions to negatively impact signal delivery. Probably more signal capacity.¹

¹ Given the same configuration.

The solution

- XCF will intermix signal sizes on “XCF Managed” signal paths as it sees fit while maintaining signal throughput and timely signal transfer, especially for small signals
 - Thus eliminating the need for transport class based signal size segregation
 - Any available path can be used for any size signal
 - Could do this today if we like. The challenge is to achieve this without incurring any discernable impact to signal throughput, system overhead, resiliency, ...

We assume adequate signal capacity and system performance.

If inadequate capacity, or systems not performing well:

- Transport classes are irrelevant; they cannot address such issues
- Segregating signal traffic on the sending system does not resolve issues related to the target system being unresponsive

Feedback during development

- On the use of existing transport class definitions:
 - Either honor them as given, or ignore them outright
 - Want to control whether system uses new or old behavior
- Should be able to observe the new behavior
 - Want to be able to see which behavior is in play, new or old
- Focus first on eliminating the need for “size segregation”
 - Applicable to most if not all installations
 - For those that don’t do group segregation, the “problem” is solved
- Don’t change meaning of existing measurements (SMF data)
 - Differences between releases makes it hard to interpret the data

← So there is a new switch, XTCSIZE

← So there is a new XCF defined pseudo transport class, _XCFMGD

Usage & Invocation

- New XCF FUNCTIONS switch XTCSIZE determines what transport class segregation rules are available to XCF
- When XTCSIZE is DISABLED:
 - XCF signal resources will be managed per traditional transport class segregation rules
- When XTCSIZE is ENABLED:
 - XCF has the option to manage signal resources for **selected** transport classes per new “XCF Managed” segregation rules
 - Transport Classes subject to being XCF Managed per XTCSIZE are those that are defined purely for size segregation

Interactions & Dependencies

- Exploitation requires:
 - A z/OS V2R4 sending system with XTCSIZE switch ENABLED
 - A target system running z/OS V2R4
 - Includes self, but the more interesting case is when target is some other system
- Software Dependencies
 - None.
- Hardware Dependencies
 - None.
- Exploiters
 - Implicit (potentially any user of XCF Signal Service)

Migration & Coexistence Considerations

- You will always see the new pseudo `_XCFMGD` “transport class”
 - But does not impact system behavior unless XTCSIZE is ENABLED.
- No toleration/coexistence APARs/PTFs are needed.
 - z/OS V2R4 manages signal resources per the traditional transport class segregation rules when communicating with a system running an older z/OS release (the XTCSIZE switch setting is irrelevant).
 - So until all systems in the sysplex are running z/OS V2R4, you will still need your traditional transport class definitions.
- Upgrade actions:
 - XTCSIZE is ENABLED by default. So ...
 - If you don't want the new behavior, XTCSIZE must be DISABLED:
 - COUPLExx parmlib member: `FUNCTIONS DISABLE(XTCSIZE)`
 - Operator command: `SETXCF FUNCTIONS,DISABLE=XTCSIZE`

Installation

- IPL z/OS V2R4
 - XTCSIZE is ENABLED by default.
 - You will get new behavior upon IPL.
- Note that XTCSIZE is a “local” switch.
 - A system obeys its own local switch setting
 - If ENABLED, there is the potential for new behavior when sending signals to any system in the sysplex that is also running z/OS V2R4
 - The XTCSIZE switch setting on some other system in the sysplex does not influence local system behavior at all.
 - Purely a “send side” setting.

Which classes are “XCF Managed” per XTCSIZE?

```
CLASSDEF CLASS(SML) CLASSLEN(956) GROUP(UNDESIG)
CLASSDEF CLASS(MED) CLASSLEN(12000)
CLASSDEF CLASS(BIG) CLASSLEN(32000)

CLASSDEF CLASS(BADGUY0) CLASSLEN(956) GROUP(MYAPP1)
CLASSDEF CLASS(BADGUY1) CLASSLEN(4028) GROUP(MYAPP1,UNDESIG)
CLASSDEF CLASS(BADGUY2) CLASSLEN(8124) GROUP(UNDESIG)
```

XTCSIZE selects those transport class definitions that do NOT have a group explicitly assigned to the class.

Which classes have explicitly assigned groups? These will NOT be XCF Managed. All others will be.

Classes BADGUY0 and BADGUY1 have an explicit group assignment, namely group MYAPP1.

GROUP(UNDESIG) is NOT an explicit group assignment. UNDESIG refers to all groups NOT explicitly assigned.

Answer: XTCSIZE would manage/select these classes:

DEFAULT – this class is implicitly defined with GROUP(UNDESIG) if you don’t otherwise change it.

SML - UNDESIG is not an explicit group assignment.

MED - no explicit groups assigned, so pure size.

BIG - no explicit groups assigned, so pure size.

BADGUY2 – UNDESIG is not an explicit group assignment.

BADGUY0 and BADGUY1 do suggest a desire for size segregation, but it’s not “pure size”. They are for group segregation.

A subtle point on class selection

```
CLASSDEF CLASS(SML) CLASSLEN(956) GROUP(UNDESIG)
CLASSDEF CLASS(MED) CLASSLEN(12000)
CLASSDEF CLASS(BIG) CLASSLEN(32000)

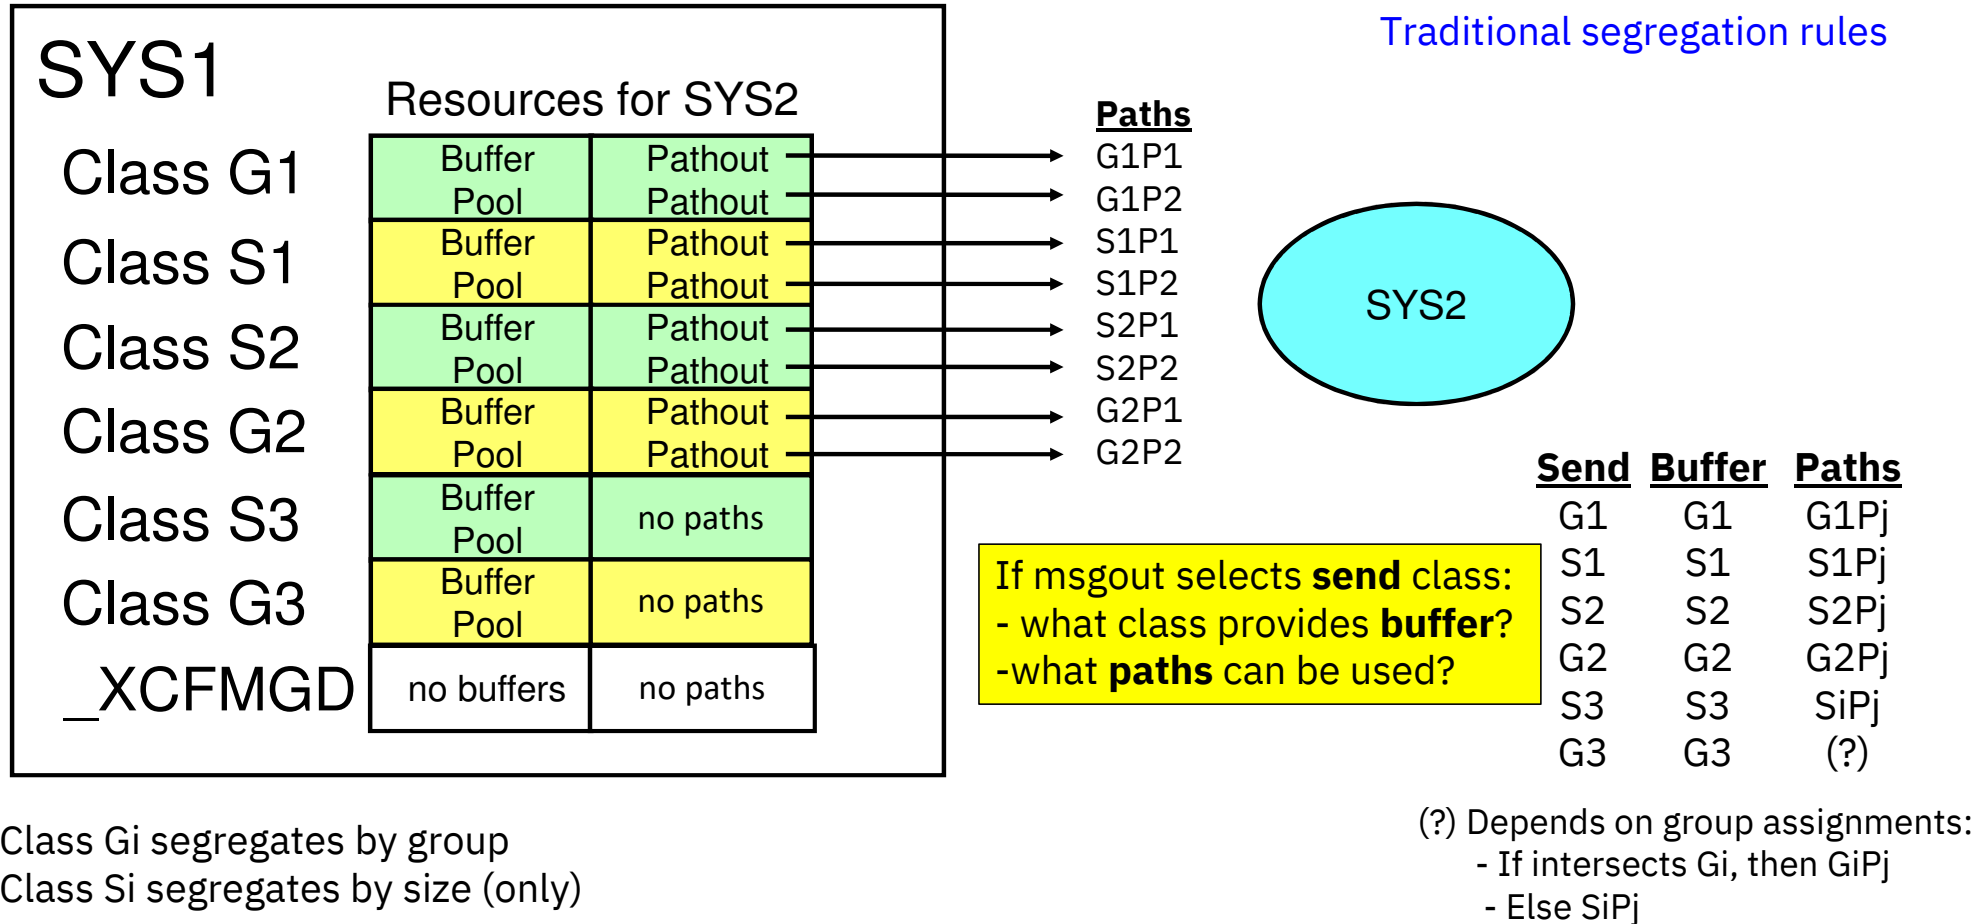
CLASSDEF CLASS(BADGUY0) CLASSLEN(956) GROUP(MYAPP1)
CLASSDEF CLASS(BADGUY1) CLASSLEN(4028) GROUP(MYAPP1,UNDESIG)
CLASSDEF CLASS(BADGUY2) CLASSLEN(8124) GROUP(UNDESIG)
```

Class BADGUY1 is actually available for use by any group, since MYAPP1+UNDESIG = everyone. Assume XTCSIZE is ENABLED.

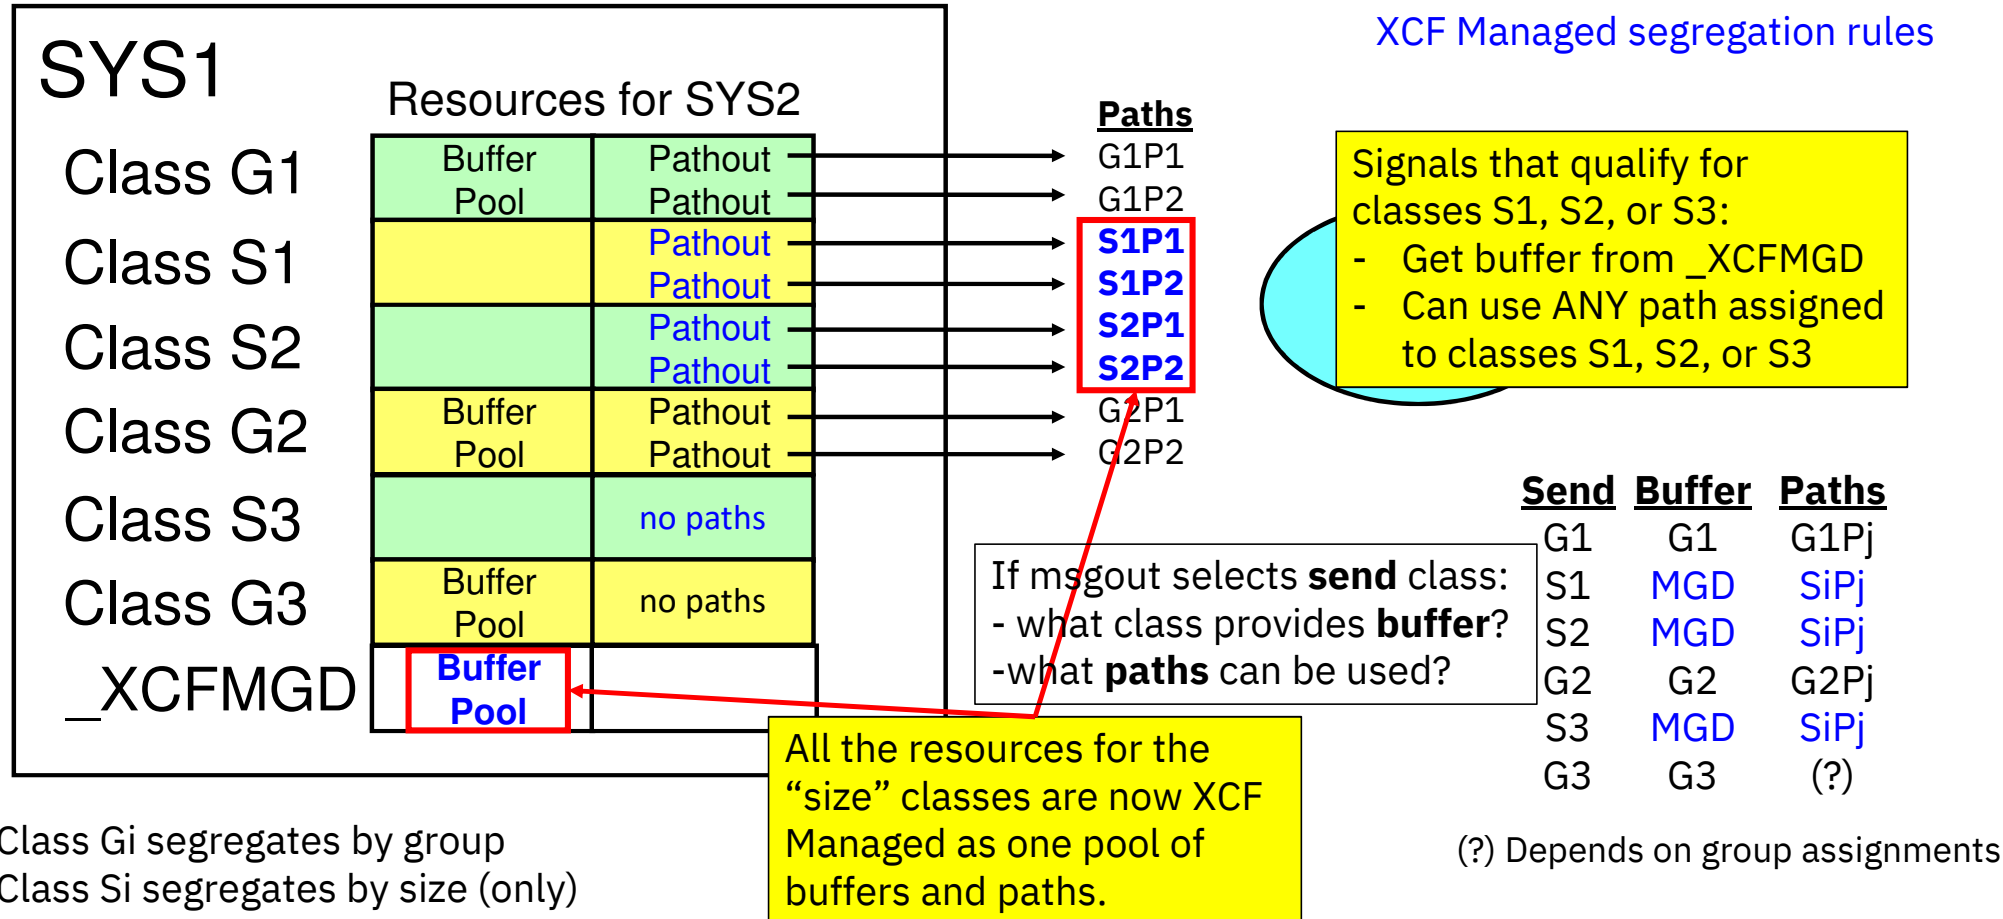
When an UNDESIG group sends a signal, both _XCFMGD and BADGUY1 are candidate classes. BADGUY1 could be selected, and likely will be if BADGUY1 has paths and CLASSLEN(4028) is a best fit.

So don't make the mistake of thinking that all UNDESIG signals must flow through _XCFMGD.

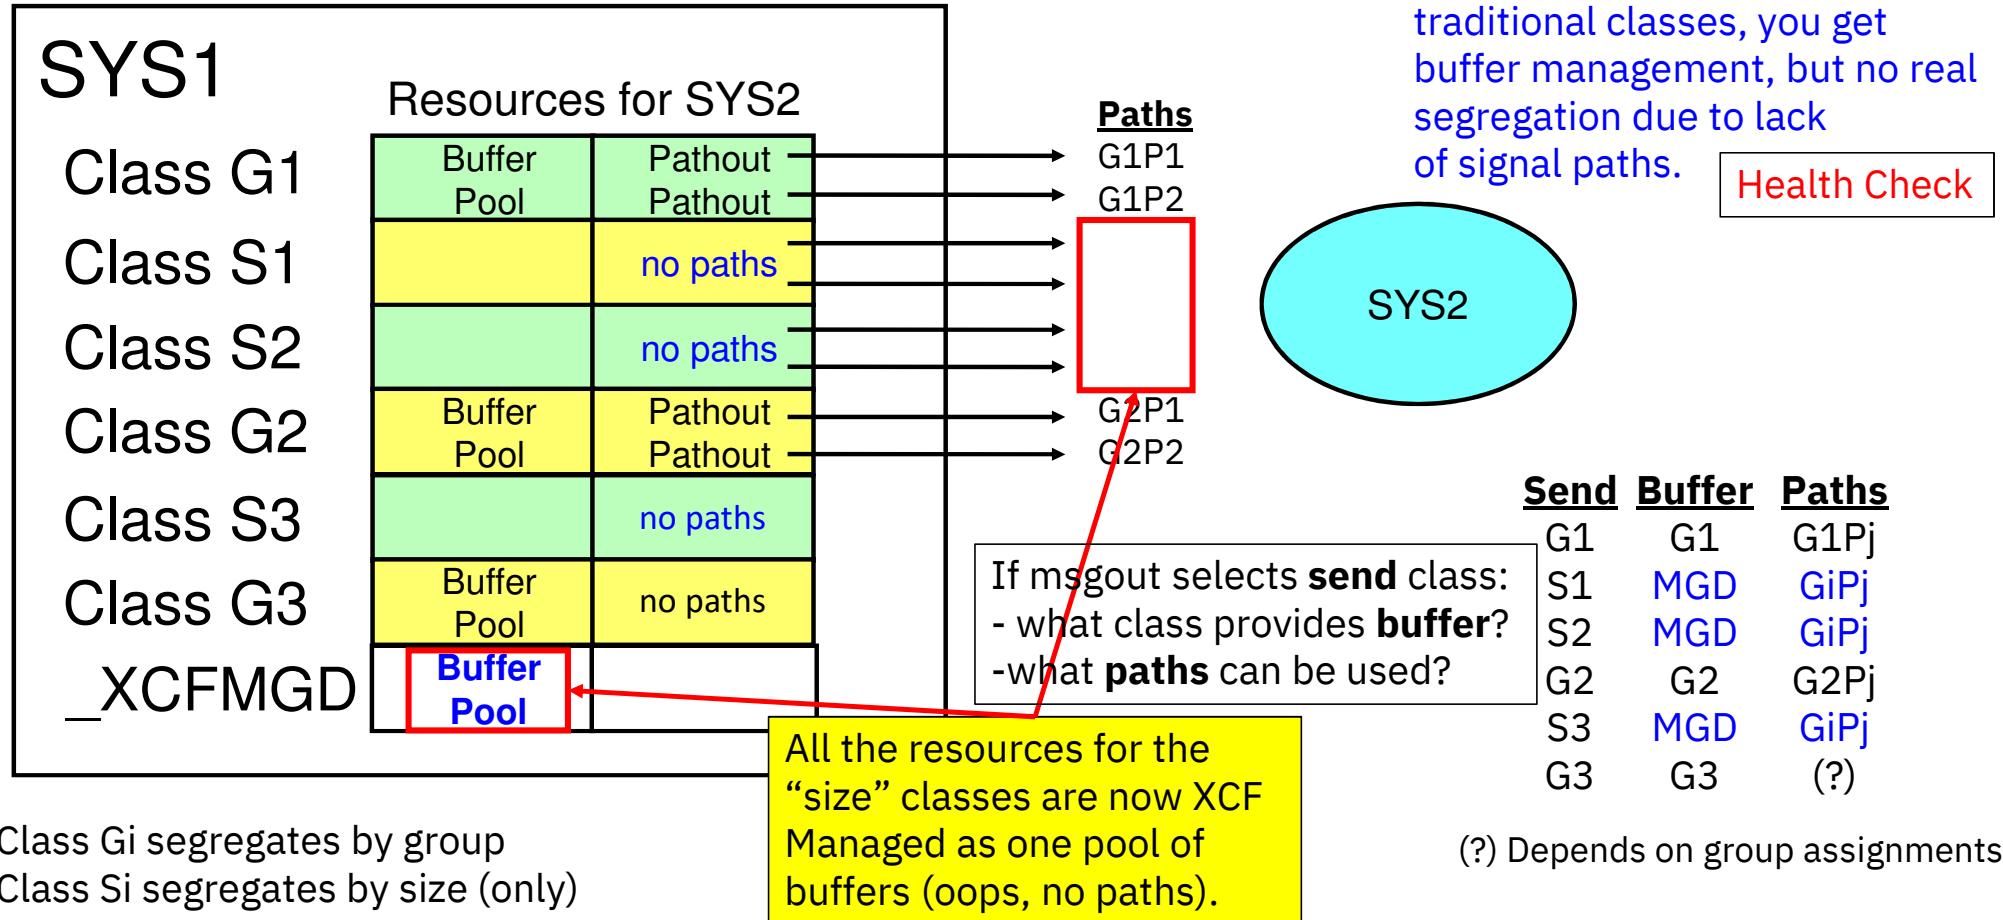
Transport Classes when XTCSIZE is DISABLED



Transport Classes when XTCSIZE is ENABLED



Transport Classes when XTCSIZE is ENABLED



About the _XCFMGD pseudo transport class

- Uses “best fit” buffers on the send side
 - Maximizes number of signals that can be accepted for a given MAXMSG limit
 - Which is important for handling bursts and delays
 - Traditional classes generally use the defined size which might not be best fit
 - So could encounter “no buffer” condition sooner than with _XCFMGD
- Paths run at the maximum signal size
 - So any size signal can be transmitted without any additional overhead
 - Never any need to re-negotiate (or tune) the paths
 - But that implies buffers on target system are likely bigger than needed
 - Which raises “no buffer” concerns (resiliency, capacity)

Oversize inbound buffers?

- You can actually sustain high signal rates with very few buffers, provided those buffers are recycled quickly
 - For a very long time, we've been doing "buffer migration" to reclaim our buffers when member message delivery takes "too long" to free them
 - As of z/OS V2R2, "Message Isolation" allows us to do "aggressive migration" to reclaim more buffers more quickly
 - With z/OS V2R4, we now have "expeditious buffer replenishment" to eliminate much of the latency that elongated "no buffer" resolution
- So we believe we can now sustain high signal rates despite the potential for having fewer inbound buffers for given MAXMSG

No significant increase in “no buffer” conditions

DISABLED

		INBOUND TO CB88		
FROM SYSTEM	T FROM/TO Y DEVICE, OR P STRUCTURE	REQ IN	BUFFERS UNAVAIL	TRANSFER TIME
CB86	S IXCPLEX_PATH1	1060K	210	0.156
	S IXCPLEX_PATH2	50	0	0.313
	S IXCPLEX_PATH3	1337K	289	0.139
	S IXCPLEX_PATH4	151	0	0.190
	C C580 TO C564	62,268	0	0.196
	C C581 TO C565	61,871	1,769	0.179
	C C582 TO C566	624	0	0.227
CB86	S IXCPLEX_PATH1	109	0	0.310
	S IXCPLEX_PATH2	3,220	93	0.280
	S IXCPLEX_PATH3	1767K	2,333	0.131
	S IXCPLEX_PATH4	1,985	0	0.143
	C C580 TO C564	5,897	0	0.173
	C C581 TO C565	85,084	0	0.179
	C C582 TO C566	70,074	209	0.202

2,521,964 REQ IN

2,268 BUFFERS UNAVAIL

1,933,369 REQ IN

2,635 BUFFERS UNAVAIL

ENABLED

Yes, “no buffer” increased.
But certainly not to the extent one might expect based on relative buffer sizes.
And not enough to impact throughput.

Oversize inbound buffers - implications

- More buffer space consumed by the normal steady state working set (likely fewer in number, but they are larger)
 - Inbound list paths will drop to zero buffers in use if there is not enough traffic to keep them busy.
 - Inbound CTC devices generally have at least 4 buffers always in use.
- But may not need as much space for the peaks
 - We will still consume up to the MAXMSG limit you provide
 - For those that have been increasing inbound MAXMSG to avoid “no buffer” conditions, you may not need such large values
 - Indeed, large values can induce long queue effects
 - In some of our tests, throughput degraded as MAXMSG increased!

Inbound “no buffer” conditions

- Our use of large buffers implies “fewer” buffers for a given MAXMSG value, which increases potential for “no buffer”
- Historically, “no buffer” seems to incent people to increase their MAXMSG values, sometimes to unreasonably high values
 - Generally, if you hit MAXMSG limit, you have queueing. The bigger the MAXMSG limit that you hit, the more queueing you have.
 - Driving “no buffer” to zero looks nice, but that queue containing tens of thousands of signals could be more detrimental to system performance than the “no buffer” condition itself
- On the inbound side, “no buffer” is not necessarily bad

Inbound “no buffer” counts

- For a CTC device, I/O is started in anticipation of there being a signal to receive. If “no buffer”, we could not start the desired I/O.
 - But if there is nothing to receive, there is no signal transfer delay.
 - So long as we get a buffer and start the I/O before a signal is sent, there is no impact no matter how high the “no buffer” count climbs
- For a list path, I/O is started when we believe there is a signal to receive. So a “no buffer” condition likely implies delay.
 - But not necessarily. Sometimes our belief is unfounded.
- XCF abhors inbound “no buffer” conditions, so it repeatedly looks to resolve them, which can drive up the “no buffer” count.
- But the count tells us very little about the impact on signal transfer since we can’t tell how much delay was induced by the “no buffer” condition.

Measuring impactful “no buffer” conditions

- We are providing additional inbound path measurements to capture the notion of an “impactful no buffer condition”
 - If we hit a “no buffer” condition, did it impact signal transfer?
 - If so, how long did the impact last?
- To some degree, we wanted insight as to whether our use of largest signal buffers (and therefore smaller number) was inducing signal transfer delays.
- But more generally, we anticipate this data will inform judgments about how many signal paths are needed:
 - Maybe we need more buffers instead of a new path
 - Maybe we need buffers to get recycled more quickly
 - Maybe we don’t care because the impacts are inconsequential

XCF defines the _XCFMGD pseudo transport class

- You cannot directly modify its attributes. However:
 - Buffer space MAXMSG limit used for _XCFMGD takes into account:
 - Default MAXMSG value
 - Buffer space limits of the classes that are being managed
 - Our own judgment as to what is a reasonable value
 - So changes to default MAXMSG value, or MAXMSG values for a managed class or a managed path could cause the buffer limits for _XCFMGD to change
 - Though you might need to DISABLE and then ENABLE the XTCSIZE switch to have them take effect
- You cannot directly assign paths to the _XCFMGD class
 - Paths are “inherited” from the classes managed when XTCSIZE is ENABLED
 - So you continue to assign paths to a traditional transport class
 - But they are reported as _XCFMGD when XTCSIZE is ENABLED
 - This lets you “see” the new behavior

DISPLAY XCF accepts _XCFMGD as a class name.
SETXCF and COUPLExx do not.

What will you see?

DISPLAY XCF output

RMF reports of XCF activity

D XCF,C – What is the XTCSIZE switch setting?

D XCF,C

IXC357I 16.08.07 DISPLAY XCF
SYSTEM SY1 DATA

....

OPTIONAL FUNCTION STATUS:

FUNCTION NAME	STATUS	DEFAULT
DUPLEXCF16	DISABLED	DISABLED
SYSSTATDETECT	ENABLED	ENABLED
USERINTERVAL	ENABLED	DISABLED
CRITICALPAGING	DISABLED	DISABLED
DUPLEXCFDIAG	DISABLED	DISABLED
CFLCRMGMT	DISABLED	DISABLED
COUPLINGTHININT	ENABLED	ENABLED
CFSTRQMON	DISABLED	DISABLED
MSGISO	ENABLED	ENABLED
XTCSIZE	ENABLED	ENABLED

....

XCF Managed is active on this system.

When sending signals to a target system running z/OS V2R4, the signal resources for all of the “size only” transport classes are being XCF Managed as a single pool that supports any size signal.

When sending signals to a down level target system, traditional transport classes and traditional segregation rules are used.

D XCF,C – What is the XTCSIZE switch setting?

D XCF,C

IXC357I 16.08.07 DISPLAY XCF
SYSTEM SY1 DATA

....

OPTIONAL FUNCTION STATUS:

FUNCTION NAME	STATUS	DEFAULT
DUPLEXCF16	DISABLED	DISABLED
SYSSTATDETECT	ENABLED	ENABLED
USERINTERVAL	ENABLED	DISABLED
CRITICALPAGING	DISABLED	DISABLED
DUPLEXCFDIAG	DISABLED	DISABLED
CFLCRMGMT	DISABLED	DISABLED
COUPLINGTHININT	ENABLED	ENABLED
CFSTRQMON	DISABLED	DISABLED
MSGISO	ENABLED	ENABLED
XTCSIZE	DISABLED	ENABLED

....

XCF Managed is not active.

You are using traditional transport
classes and traditional transport
class segregation rules.

Change XTCSIZE switch setting - dynamically

SETXCF FUNCTIONS,DISABLE=XTCSIZE

Revert to old behavior

IXC373I XCF / XES OPTIONAL FUNCTIONS DISABLED:
XTCSIZE

SETXCF FUNCTIONS,ENABLE=XTCSIZE

Enable new behavior

IXC373I XCF / XES OPTIONAL FUNCTIONS ENABLED:
XTCSIZE

Change XTCSIZE switch setting – through IPL

COUPLExx parmlib member

COUPLE SYSPLEX(plexname)
 PCOUPLE(...)
 ACOUPLE(...)

FUNCTIONS ENABLE(XTCSIZE)

CLASSDEF ... Enable new behavior
PATHIN ...
PATHOUT ...

COUPLExx parmlib member

COUPLE SYSPLEX(plexname)
 PCOUPLE(...)
 ACOUPLE(...)

FUNCTIONS DISABLE(XTCSIZE)

CLASSDEF ... Revert to old behavior
PATHIN ...
PATHOUT ...

D XCF,CLASSDEF – to see transport class data

```
D XCF,CLASSDEF,CLASS=ALL
```

```
IXC344I 16.08.23 DISPLAY XCF 848
```

TRANSPORT	CLASS	DEFAULT	ASSIGNED
CLASS	LENGTH	MAXMSG	GROUPS
_XCFMGD	0	2000	UNDESIG
BADGUY0	956	2000	MYAPP1
BADGUY1	4028	2000	MYAPP1
BADGUY2	8124	2000	UNDESIG
BIG	32000	2000	UNDESIG
DEFAULT	956	2000	UNDESIG
MED	12000	2000	UNDESIG
SML	956	2000	UNDESIG

Always
present →

UNDESIG

These classes
selected when
XTCSIZE is
ENABLED.

We don't have anything to show you which classes qualify for XTCSIZE management.

D XCF,CLASSDEF – to see transport class data (cont)

As for any traditional class, _XCFMGD signal size distributions are shown for each target system.

....

_XCFMGD TRANSPORT CLASS USAGE FOR SYSTEM SY1

SUM MAXMSG:	10000	IN USE:	2	NOBUFF:	0
SEND CNT:	978	BUFFLEN (FIT):	956		
SEND CNT:	362	BUFFLEN (BIG):	4028		
SEND CNT:	21	BUFFLEN (BIG):	8124		

SY1 is local system

_XCFMGD TRANSPORT CLASS USAGE FOR SYSTEM SY2


SUM MAXMSG:	22000	IN USE:	8	NOBUFF:	0
SEND CNT:	1713	BUFFLEN (FIT):	956		
SEND CNT:	130	BUFFLEN (BIG):	4028		
SEND CNT:	119	BUFFLEN (BIG):	8124		

SY2 running z/OS V2R4
("XCF Managed" is in play)

_XCFMGD TRANSPORT CLASS USAGE FOR SYSTEM SY3

SUM MAXMSG:	16000	IN USE:	0	NOBUFF:	0
SEND CNT:	0	BUFFLEN (FIT):	956		

SY3 running z/OS V2R3
Down level, so old rules.
("XCF Managed" not used)



Which list paths are XCF Managed ?

D XCF,PO,STRNAME=ALL,STATUS=WORKING

IXC356I 16.09.28 DISPLAY XCF 855

STRNAME	REMOTE	PATHOUT	UNUSED		TRANSPORT
PATHOUT	SYSTEM	STATUS	PATHS	RETRY	MAXMSG CLASS
IXCTL_SIGNAL01		WORKING	234	10	2000 DEFAULT
	SY2	WORKING			<u>_XCFMGD</u>
	SY3	WORKING			DEFAULT
IXCTL_SIGNAL02		WORKING	234	10	2000 DEFAULT
	SY2	WORKING			_XCFMGD
	SY3	WORKING			DEFAULT

Structure still indicates the "home" class

Uplevel target system

Downlevel target system (so using old rules)

STRNAME	REMOTE	PATHOUT	TRANSFR	BUFFER	MSGBUF	SIGNL	MXFER
PATHOUT LIST	SYSTEM	STATUS	PENDING	LENGTH	IN USE	NUMBR	TIME
IXCTL_SIGNAL01							
	12 SY2	WORKING	0	<u>62464</u>	204	10	2915
	9 SY3	WORKING	0	956	6	3	4612
IXCTL_SIGNAL02							
	12 SY2	WORKING	0	62464	270	11	2767
	9 SY3	WORKING	0	956	8	4	1049

Running at biggest size

Running at defined size

Same list paths after XTCSIZE is DISABLED

D XCF,PO,STRNAME=ALL,STATUS=WORKING

IXC356I 16.10.32 DISPLAY XCF 871

No class is XCF Managed

STRNAME	REMOTE	PATHOUT	UNUSED	TRANSPORT
PATHOUT	SYSTEM	STATUS	PATHS	RETRY MAXMSG CLASS
IXCTL_SIGNAL01		WORKING	234	10 2000 DEFAULT
	SY2	WORKING		DEFAULT
	SY3	WORKING		DEFAULT
IXCTL_SIGNAL02		WORKING	234	10 2000 DEFAULT
	SY2	WORKING		DEFAULT
	SY3	WORKING		DEFAULT

Reverts to "home" class

STRNAME	REMOTE	PATHOUT	TRANSFR	BUFFER	MSGBUF	SIGNL	MXFER
PATHOUT LIST	SYSTEM	STATUS	PENDING	LENGTH	IN USE	NUMBR	TIME
IXCTL_SIGNAL01							
12 SY2	WORKING	0	956	12	18	3418	
9 SY3	WORKING	0	956	8	4	3044	
IXCTL_SIGNAL02							
12 SY2	WORKING	0	956	14	19	262	
9 SY3	WORKING	0	956	10	5	3116	

Reverts to defined size

Which CTC signal paths are XCF Managed?

D XCF,PO,DEV=ALL,STATUS=WORKING
IXC356I 16.09.14 DISPLAY XCF 852

LOCAL DEVICE	REMOTE	PATHOUT	REMOTE	TRANSPORT
PATHOUT	SYSTEM	STATUS	PATHIN	RETRY
8000	SY2	WORKING	80BF	10
8001	SY2	WORKING	80BE	10
8002	SY3	WORKING	80BD	10
8003	SY3	WORKING	80BC	10

DEFAULT class is XCF Managed.
BADGUY0 is not.

You can't see the
"home" class

Uplevel target system

Downlevel target system
(so using old rules)

LOCAL	REMOTE	REMOTE	PATHOUT	TRANSFR	BUFFER	MSGBUF	SIGNL	MXFER
PATHOUT	PATHIN	SYSTEM	STATUS	PENDING	LENGTH	IN USE	NUMBR	TIME
8000	80BF	SY2	WORKING	0	956	14	826	34
8001	80BE	SY2	WORKING	0	62464	14	1387	38
8002	80BD	SY3	WORKING	0	956	14	1193	48
8003	80BC	SY3	WORKING	0	956	14	1847	29

Running at biggest size

Running at defined size

Same CTC paths after XTCSIZE is DISABLED

No class is XCF Managed

```
D XCF,PO,DEV=ALL,STATUS=WORKING
IXC356I 16.10.28 DISPLAY XCF 868
```

LOCAL DEVICE	REMOTE	PATHOUT	REMOTE		TRANSPORT
PATHOUT	SYSTEM	STATUS	PATHIN	RETRY	MAXMSG CLASS
8000	SY2	WORKING	80BF	10	2000 BADGUY0
8001	SY2	WORKING	80BE	10	2000 DEFAULT
8002	SY3	WORKING	80BD	10	2000 BADGUY0
8003	SY3	WORKING	80BC	10	2000 DEFAULT

Reverts to "home" class

LOCAL	REMOTE	REMOTE	PATHOUT	TRANSFR	BUFFER	MSGBUF	SIGNL	MXFER
PATHOUT	PATHIN	SYSTEM	STATUS	PENDING	LENGTH	IN USE	NUMBR	TIME
8000	80BF	SY2	WORKING	0	956	14	828	35
8001	80BE	SY2	WORKING	0	956	14	1497	36
8002	80BD	SY3	WORKING	0	956	14	1194	29
8003	80BC	SY3	WORKING	0	956	14	1950	23

Reverts to defined size

RMF Report: XCF Usage by System

OUTBOUND FROM S5A									
TO SYSTEM	TRANSPORT CLASS	BUFFER LENGTH	REQ OUT	% SML	% FIT	% BIG	% OVR	ALL PATHS UNAVAIL	REQ REJECT
S5B	_XCFMGD	956	624,950	0	99	1	100	0	0
	BIG	40,892	0	0	0	0	0	0	0
	CTTX	40,892	0	0	0	0	0	0	0
	DAE	956	0	0	0	0	0	0	0
	DEFAULT	20,412	0	0	0	0	0	0	0
	DEFSMALL	956	0	0	0	0	0	0	0
	DEF8K	8,124	0	0	0	0	0	0	0
	FEWFAST	956	0	0	0	0	0	0	0

_XCFMGD transport class will always appear.

- Its counts will be zero if XTCSIZE is DISABLED or if the target system is down level.

- SML, FIT, BIG are truthful and consistent with existing definitions of SMF data.

In the past, these numbers might have prompted one to consider revising class definitions.

But with _XCFMGD, not only is there nothing to be done, there is nothing you can do.

The traditional classes will always appear as well. The counts will be zero if the class is not being used.

But from this report, you can't tell why there's no use: XCF Managed? No signals sent? (CTTX for example).

RMF Report: XCF Path Statistics

OUTBOUND FROM S5A								
TO SYSTEM	T Y	FROM/TO DEVICE, OR STRUCTURE	TRANSPORT CLASS	REQ OUT	AVG Q LENGTH	AVAIL	BUSY	RETRY
S5B	S	IXCPLEX_PATH1	_XCFMGD	147,251	0.00	147,251	0	0
	S	IXCPLEX_PATH2	_XCFMGD	454,322	0.00	454,322	0	0
	S	IXCPLEX_PATH3	_XCFMGD	5,011	0.00	5,011	0	0
	S	IXCPLEX_PATH4	CTTX	0	0.00	0	0	0
	S	IXCPLEX_PATH5	_XCFMGD	18,395	0.00	18,383	12	0
	C	C5B0 TO C5A4	_XCFMGD	57	0.00	53	4	0
	C	C5B1 TO C5A5	_XCFMGD	73	0.00	64	9	0
	C	C5B2 TO C5A6	_XCFMGD	52	0.00	47	5	0



Paths are reported as being assigned to _XCFMGD transport class if they are being XCF Managed.

What might you see

- Most of you likely have:
 - Multiple “pure size” transport classes defined
 - One or more signal paths assigned to each of those classes
 - Very high AVAIL percentages for your paths
 - Signal paths nearly always “not busy” when picked to send a signal
 - Some paths hardly ever picked, so lots of excess capacity
 - Hardly any “no buffer” conditions
- So when XTCSIZE is ENABLED
 - Best fit buffers for _XCFMGD likely implies more send side buffer capacity than your definitions provide
 - Signals likely distributed to paths much as your traditional classes do today
 - An underutilized path in one of your classes can now be used to help a peer class that otherwise would have used a “busy” path

Signal distributions might change

•

OUTBOUND FROM CB86

TO SYSTEM	T FROM/TO Y DEVICE, OR P STRUCTURE	TRANSPORT CLASS	REQ OUT	AVG Q LENGTH	AVAIL	BUSY	RETRY
DISABLED	CB8C S IXCPLEX_PATH1	DEFSMALL	1150910	0.00	1084034	66,876	0
	S IXCPLEX_PATH2	DEFMED	931	0.00	931	0	0
	S IXCPLEX_PATH3	DEFSMALL	3301531	0.00	3225472	76,059	0
	S IXCPLEX_PATH4	DEFMED	30	0.00	30	0	0
ENABLED	CB8C S IXCPLEX_PATH1	_XCFMGD	925,035	0.00	924,888	147	0
	S IXCPLEX_PATH2	_XCFMGD	362,281	0.00	362,146	135	0
	S IXCPLEX_PATH3	_XCFMGD	3362504	0.00	3362334	170	0
	S IXCPLEX_PATH4	_XCFMGD	13,133	0.00	12,901	232	0

With XTCSIZE ENABLED

- Signals more evenly distributed across paths
- Number of “busy” becomes negligible.
- (not shown) Inbound went from 0 to 6 no-buffer conditions.

Expectations

- Overall, your sysplex workloads should run at least as well as they do now. Signal performance should not be degraded.
 - Regardless of what z/OS release the target system is running
 - Regardless of whether XTCSIZE is ENABLED or DISABLED
- For a given set of signal paths, there is no set of (pure size) transport class definitions for those paths that will perform better with XTCSIZE=DISABLED than with XTCSIZE=ENABLED.
- The only determining factors for signal performance will be:
 - The number of signal paths you provide
 - The performance characteristics of those paths
 - The performance characteristics of the systems using those paths

Possible exceptions?

- Our overall improvements to signal delivery might change timings that induce secondary impacts
- Signal traffic might flow over different paths than in the past, which might impact things like CF, subchannel, and link utilization
 - If you have classes with “busy” paths today, those signals are likely to be sent via an available path from a peer XCF Managed class.
 - So maybe your high frequency signals always traveled via CF1. With XTCSIZE=ENABLED, any signal can go via any XCF Managed path. So CF2, which in the past got very little activity, might now be used more frequently.
- Use of max size inbound buffers might be an issue for installations that are memory constrained and/or have lots of signal paths (especially CTC devices)

Session Summary

- We have eliminated the need for you to define XCF Transport Classes to segregate signals by size
 - No more planning, defining, monitoring, tuning, changing, ...
- Just ENABLE the XTCSIZE switch. We'll do the rest.
 - In particular, you don't need to change your current XCF transport class configuration

Appendix

- Contacts
 - Mark A Brooks – mabrook@us.ibm.com
 - Neil Johnson - najohnsn@us.ibm.com